

Domain

mullover.ai registered through 2028

Author

Research and synthesis May 2026



BUSINESS CASE

Edition Two

REVISED · MAY 2026

Contents

PART I • THE PRODUCT

1	Executive summary	3
2	What Mull Over actually does, in plain English	5
3	The killer feature: divergence analysis	8
4	The eight question formats, including photos	11
5	Use cases written like you are twelve	13
6	Mull Over for Teams: the corporate use case	17

PART II • THE MARKET AND THE COMPETITION

7	Market sizing and the realistic addressable slice	20
8	Direct competitors in couples technology	22
9	Couples therapy apps and clinical platforms	26
10	Decision-making and group polling tools	29
11	Photo voting and image polling tools	32
12	Feature matrix and where Mull Over wins	34
13	Why the gap is real and has stayed open	36

PART III • THE BUSINESS

14	Pricing model for consumer and teams	38
15	Unit economics, decomposed	40
16	Best initial markets and launch sequence	42

PART IV • MARKETING AND GROWTH

17	Channel overview and CAC targets	44
18	Meta Ads playbook: techniques and creative	45
19	SEO and content compounding	49
20	Therapist partnerships	50
21	PR, podcasts, and the rest	51

PART V • REVENUE, RISKS, AND EXPANSION

22	Three revenue scenarios, realistic	52
23	Probability of reaching one hundred thousand per year	54
24	Expansion paths and adjacent products	56
25	Risks and mitigations	58

PART VI • BUILD

26	Tech stack and architecture	60
27	Build plan and time to launch	62
28	Recommendation	63

SECTION 1

Executive summary

Mull Over is a thirty-day decision aggregation product. Two people who share a big decision (move, baby, job, end or stay, buy this house) each tap one button per day. They cannot see each other's votes. On day thirty the verdict reveals, with an AI synthesis that names the actual disagreement underneath the surface position. It is the only product on the market that combines private daily voting with longitudinal capture and AI-driven divergence analysis.

The market is real but the founder should not over-claim

The couples technology market is genuinely large and growing. Industry analysts size it at **around two billion US dollars in 2024, projected to reach five point seven seven billion by 2033 at a 12.5% compound annual growth rate** (Business Research Insights, 2025). The couples counseling app sub-segment is forecast to climb from seven hundred and twenty million in 2025 to two point one billion by 2032 at 14.2% CAGR (HTF Market Intelligence, 2026).

Mull Over does not compete head-on with most of this market. Its differentiation sits in a small specific corner: private longitudinal voting plus AI synthesis applied to high-stakes life decisions. The realistic addressable slice in the first three years is small but high-margin: roughly two to ten thousand paying customers across five English-speaking markets, generating between thirty thousand and four hundred thousand dollars in annual revenue depending on which growth channels work and how well the synthesis lands.

The unit economics carry the case

The price points are calibrated against established benchmarks: nineteen ninety-nine US dollars per shared decision (one-off, paid at verdict), or ninety-nine dollars per year for unlimited. Per-transaction margin is approximately ninety-five percent: ninety cents of Stripe fees plus three cents of Claude API on a twenty-dollar charge. Even modest paid volume produces meaningful contribution.

95.4%

GROSS MARGIN PER DECISION

\$19.07

NET CONTRIBUTION PER SALE

\$2,800

YEAR-ONE TOTAL INFRASTRUCTURE

The competitive landscape is more crowded than first inspection suggests

Section 8 expands the competitor analysis from the eight names in the first edition to over thirty named products across four categories: couples technology (Paired, Lasting, Relish, Coral, Between, OurCouple, OurRitual, Coupleness, Agapé, Love Nudge, Connected, Lovewick, Mend, Cupla, Raft, Honeydue, Donkey Chats), couples therapy and clinical platforms (BetterHelp, Talkspace, ReGain, Calmerry, Amwell, MDLive, Teladoc, Thriveworks, Growing Self, Couples Therapy Inc., Online-Therapy.com), decision-making and group polling (Loomio, Cloverpop, Polly, Doodle, 1000minds, Decisive, ChoiceMap), and image polling (PollUnit, Opinion Stage, Vote Social, Woobox, Traxoft).

None of these products combines Mull Over's three structural ingredients: time-series private voting, AI divergence synthesis, and per-decision pricing on a couple-specific use case. The detailed feature matrix in section 12 confirms this. The point is not that no one is competing for the customer's attention. Everyone is. The point is that nobody is offering this specific product, and the structural reasons for that gap (covered in section 13) are durable.

What this document is and is not

This document is a tool for the founder's own decision-making. It is intentionally honest about the risks and the realistic scale. It is not a fundraising deck and is not optimised for investor signalling. The recommendation in section 28 still favours building, but on a self-funded, low-overhead, conservative basis with a clear exit path if the synthesis quality fails to land with the first one hundred paying customers.

I

The product, the killer

SECTION 2

feature and who it helps

What Mull Over actually does, in plain English

Two people who share a decision each tap one button per day for thirty days. They cannot see each other's votes during that period. On day thirty the verdict reveals both votes, side by side, plus an AI-generated synthesis that names the actual disagreement underneath. Then they have the conversation with having English language games for couples and beyond, and the corporate version for teams.

The three steps

STEP ONE: POSE THE QUESTION

The user enters the decision in their own words. "Should we move to Wellington?" "Should we have a baby in 2026?" "Should I leave my job?" "Should we buy this house?" They pick a vote format (more on the eight formats in section 4), set a duration (default thirty days, range one to three hundred and sixty-five), and choose whether to mull solo or invite a partner. If they invite a partner, the partner receives a magic-link email and joins by signing in. No passwords.

STEP TWO: VOTE EACH DAY

Once per day, both partners get a notification (push to phone, email if push is unavailable). They tap one of the available options (Strong No, Lean No, Lean Yes, Strong Yes for the default format) and optionally write a one-line note. "Walked through town at lunch, felt great." "Spent the afternoon in the garden, would hate to lose this." Total time per vote: fifteen seconds. Total time per partner across thirty days: about seven minutes.

STEP THREE: THE VERDICT REVEALS

On day thirty, both partners get an email and a push. They open the app and see, for the first time, the other person's complete history of votes laid alongside their own. James voted GO with a positive twelve weighted score. Alexandra voted LEAN NO with a negative three. Below the verdicts is a day-by-day chart, then the AI synthesis paragraph, then the conversation prompt.

The privacy mechanic is non-negotiable

While the thirty days run, neither partner can see how the other is voting. They can see whether the other has voted today, never which direction. Notes are sealed until day thirty. This is the central design rule of the product and it is not optional, not a settings toggle, not a premium override.

The reason is structural. Without privacy, whoever votes first sets the anchor. Whoever votes loudest pulls the mood. Both partners drift toward each other to avoid friction. The signal becomes a negotiated agreement rather than two independent perspectives. With privacy, each partner is talking to the app, not to each other. They can register doubt or excitement that would feel embarrassing to say out loud. On day thirty, the verdict reveals the real positions, and that becomes the foundation for the conversation.

DESIGN RULE, WRITTEN INTO THE SPEC

Votes are sealed until verdict. No exceptions, no settings toggle, no premium override. If a couple wants to share daily votes with each other, they can do that without an app. Mull Over exists specifically because that approach does not work.

What the AI synthesis actually says

The synthesis paragraph is the centrepiece of the verdict screen. It reads the daily notes from both partners, deduplicates and clusters the themes, counts frequencies, and writes a short paragraph that names the disagreement underneath the surface vote. For the "should we move to town" example with James and Alexandra:

SAMPLE SYNTHESIS OUTPUT

"You are not actually split on whether to move. You are split on what you would give up. James kept returning to walkability, cafes, and social life across twelve different notes. Alexandra kept returning to the garden, the quiet, and the outdoor space you have built across fourteen notes. Your strongest disagreements clustered in week three, after a sequence of open-home weekends."

And below the synthesis, the conversation prompt:

SAMPLE CONVERSATION PROMPT

"Can you find a place in town that protects what you would both miss most: real outdoor space and quiet enough to breathe?"

That second box is the most useful single output of the entire product. It turns an abstract disagreement (move or stay) into a concrete design problem the couple can solve together (find a third option that addresses both).

What Mull Over is not

NOT THERAPY

Therapy is a guided conversation. Mull Over comes before the conversation, producing the structured data that makes the conversation more honest. Many therapists interviewed during validation said they would happily refer clients to a tool that does this preparatory work because they spend the first three sessions of every new couple doing it manually.

NOT A DAILY-QUESTION APP

Paired and Lasting send daily questions to spark conversation ("what is something your partner did this week that meant a lot to you"). That is relationship maintenance. Mull Over is decision support. The two products serve different needs and a couple might use both.

NOT A COACHING SERVICE

Relish, OurRitual, and others charge thirty-two to ninety-six dollars per week for human coach access. Mull Over is software, deliberately. The unit economics that make ninety-nine dollars per year viable depend on no human in the loop. The AI synthesis substitutes for a coach reading the journal and reflecting back the patterns. It is not as deep as a coach, but it is roughly thirty times cheaper and entirely private.

NOT A POLLING TOOL

SurveyMonkey, Google Forms, Doodle, Polly, and similar do single-event polls. They are point-in-time captures. Mull Over is a longitudinal capture over thirty days, which is the structural difference that lets the synthesis surface trajectory and pattern rather than just current opinion.

Free versus paid

The free tier is deliberately generous, not a stub. One active solo decision, all eight vote formats, durations up to ninety days, edit and delete and pause and resume, anonymous public share link, basic CSV export, and PWA install. The deterministic verdict (the math part, without the AI synthesis) is included free. This is the funnel: a user who runs a free solo decision and gets a clean well-designed verdict on day thirty is primed to bring their partner in for the next decision.

The paid tiers unlock two things, both genuinely worth paying for: **partner mode** (the killer feature, since divergence analysis is impossible without two sides of data) and **AI insight** (the synthesis, theme extraction, conversation prompt). Per-decision pricing at nineteen ninety-nine lets a couple pay for one specific big decision without subscribing. Annual at ninety-nine dollars is for users who will mull more than five decisions in a year.

SECTION 3

The killer feature: divergence analysis

Divergence analysis is the single output that justifies the nineteen ninety-nine per-decision price. It is what people are actually paying for, and the feature that makes Mull Over genuinely defensible. Five layers, half deterministic, half AI.

The five layers

LAYER ONE: AGREEMENT RATE

For every day in the period, the system checks whether both partners leaned in the same direction. The output is a single percentage: "you agreed on sixteen of thirty days" or "you diverged on forty-seven percent of days." This is the simplest measurement and it carries weight because the size of the agreement gap is itself informative. Two partners who agreed every day but disagreed on the final tally have a different problem from two who agreed barely half the time.

LAYER TWO: CONVICTION TRAJECTORY

A regression line is fit through each partner's daily votes (weighted negative two to positive two for the Strong/Lean format). The slope is reported as "positive zero point four per week" or "negative zero point two per week." Positive slope means conviction is rising. Negative means conviction is softening. The divergence in trajectories tells you who is growing into the decision and who is growing out of it.

LAYER THREE: THEME EXTRACTION (AI)

The first AI-powered layer. Claude reads every note each partner wrote across the thirty days, deduplicates, clusters into themes, and counts frequencies. Output: "James mentioned walkability twelve times. Alexandra mentioned garden and outdoor space fourteen times." The frequency counts make the themes feel grounded, not invented. Couples can verify them against their own memory of the period.

LAYER FOUR: THE SYNTHESIS PARAGRAPH (AI)

The most valuable single output. Claude takes both theme breakdowns and writes a short paragraph naming the actual disagreement underneath the surface vote. The structure is consistent: "You are not actually split on X. You are split on Y." The re-framing is what people remember and quote back to their friends.

LAYER FIVE: THE CONVERSATION WORTH HAVING (AI)

The most actionable output. Claude generates a single specific question for the couple to take into their next real conversation. Not generic advice ("communicate openly"), but a question that addresses the specific gap surfaced in layer four. For the move example: "Can you find a place in town that protects what you would both miss most?"

WHY THESE FIVE LAYERS AND NOT THREE OR SEVEN

Five is the minimum that delivers a complete picture: two layers of structured data, two layers of synthesis, one layer of action. Fewer and the report feels thin. More (additional charts, sentiment analysis, prediction) and the user gets overwhelmed and the cost per analysis rises. Five is the deliberate ceiling.

What it costs to produce

The economics are what make nineteen ninety-nine viable. Layers one and two are free; the math runs in milliseconds. Layers three to five use a single Claude API call with a structured prompt. Verified pricing for Sonnet 4.6 (Anthropic, May 2026): three dollars per million input tokens, fifteen dollars per million output tokens.

A typical analysis for a thirty-day decision averages four thousand input tokens (system prompt plus both vote streams plus both note sets) and twelve hundred output tokens. At full price that is one point two cents input plus one point eight cents output, equals roughly three cents per analysis. With prompt caching (ninety percent off cached system prompt) and message batch pricing (fifty percent off for non-real-time), the cost falls to roughly one point five cents per analysis at scale.

PER-DECISION UNIT ECONOMICS

\$19.07 net

Nineteen ninety-nine charged, minus eighty-eight cents Stripe fees (US domestic), minus three cents Claude API, equals nineteen oh seven net contribution. Gross margin: ninety-five point four percent. At ten thousand verdicts per month, the API spend is three hundred dollars against two hundred thousand in revenue.

Why the synthesis works (when it does)

Most AI-generated text in consumer products feels generic because it is untethered from specific data. The standard failure mode of AI relationship advice is platitudes: "communication is key," "make time for each other." Couples spot this instantly and discount the entire output.

Mull Over's synthesis works because it is grounded in two specific data sources: the daily votes (which provide a structural skeleton: sixteen of thirty days agreed, James's conviction rose, Alexandra's softened) and the daily notes (which provide the human texture: walkability, garden, mortgage stretch, social life). The prompt instructs Claude to use both: name the themes from the notes, but anchor them to the trajectory data. Output sounds specific because it is specific.

Where it can fail, and the mitigations

FAILURE ONE: THEMES ARE NOT GROUNDED

If the notes are sparse (couples who skipped writing notes), the theme extraction has nothing to work with. Mitigation: the verdict screen detects this and falls back to a "low-data" version of the report that is honest about the limitation, plus a nudge to be more verbose in future decisions.

FAILURE TWO: SYNTHESIS NAMES THE WRONG DISAGREEMENT

The model gets the wrong layer of the conflict. Mitigation: prompt explicitly asks for "the underlying axis, not the surface position" and includes few-shot examples of good versus bad synthesis. During beta, every analysis is reviewed by a human before send, with the rejection reasons feeding back into prompt iteration.

FAILURE THREE: CONVERSATION PROMPT SOUNDS THERAPY-ISH

"Talk about your feelings" type platitudes. Mitigation: prompt restricts to action-oriented questions ("Can you find," "What would need to be true"). Outputs are scored against this rubric in offline evals.

All three failure modes are addressable through prompt engineering and offline evaluation before the cost ever touches a paying customer. The first one hundred paid analyses go through human review with prompt iteration. From analysis one hundred and one onward, automated evaluations catch most regressions.

Why this layer-five output is the most quoted

Across early prototype testing, users consistently quoted the conversation prompt rather than the synthesis when describing the experience to friends. "It told us to ask if we could find a place that protects what we would both miss" travels better than "it explained that we were split on what we would give up." The verdict is the headline; the conversation prompt is what people actually use. This shapes the marketing: lead with the conversation prompt examples, not the verdict graphics.

The eight question formats, including photos

Different decisions deserve different vote types. Mull Over offers all eight, all free, all at every tier. The first seven were in the original product spec. The eighth (photo voting) is a new addition for edition two, and unlocks a meaningful set of additional use cases including outfit choices, design selections, and visual A-versus-B decisions of all kinds.

The eight formats

FORMAT	BEST FOR	EXAMPLE QUESTION
Yes / No	Clean binary decisions	"Should we sell the house?"
Scale 1 to 10	How good does it feel today	"How excited are you about moving to Sydney?"
A vs B	Two named alternatives	"Wellington or Auckland?"
Strong / Lean	Four-point conviction (default for life decisions)	Strong No, Lean No, Lean Yes, Strong Yes
Pros vs Cons	Accumulating tallies	"Add one pro or one con each day"
Pick best	One of N named alternatives	"Which neighbourhood feels right?"
Rank options	Order N alternatives	"Rank these five baby names"
Photo A vs B	Visual choice between two images	"Which dress for the wedding?"

The photo format, in more detail

This is the most significant new addition for edition two. When a user picks Photo A vs B, they upload two images at decision-creation time. The two photos are stored privately in object storage (Backblaze B2 with signed URLs, encrypted at rest). Each day during the period, both partners see the two images side by side and pick one. They can also write a one-line note explaining their pick.

The daily-photo format has a special property the text formats do not: **the image content itself becomes part of the synthesis input**. Claude can be given a structured description of each image (generated at upload time by a quick vision model call, cached) so the synthesis can reference what is actually in the pictures. "James consistently picked the navy dress (image B) when his note mentioned 'formal' or 'wedding'. Alexandra picked the green dress (image A) when her note mentioned 'photos' or 'colour'."

What photo voting unlocks

OUTFIT AND STYLING DECISIONS

"Which dress for the wedding?" "Which suit for the interview?" "Which outfit for the photo shoot?" These are decisions where one person is the chooser, and the partner's input is genuinely useful. Lower-stakes than "should we move," shorter periods (often 3 to 7 days, sometimes 1 day), but high willingness to pay for a clean answer.

RENOVATION AND HOME DESIGN

"Which kitchen splashback?" "Which paint colour for the lounge?" "Which front-door colour?" Couples renovating their house spend weeks deliberating on these. Mull Over with photo format gives them a structured way to actually decide, with the AI synthesis describing which factors each partner kept weighting.

WEDDING AND EVENT PLANNING

"Which venue?" "Which bouquet?" "Which save-the-date design?" "Which invitation layout?" Weddings produce dozens of these visual A-versus-B choices, often under time pressure. A single annual subscription paid by the engaged couple covers six months of these decisions.

BABY AND PARENTING DECISIONS

"Which cot?" "Which pram?" "Which name?" (with photos as a memory aid, often paired with name-association images). New parents face an avalanche of visual decisions in their first year.

DESIGN AND CREATIVE WORK

"Which logo for our business?" "Which photo for the LinkedIn profile?" "Which Instagram post should I publish?" Solo creators and small partnerships use the photo format to get structured input rather than the friend-group chat-poll method that produces dominant-voice answers.

A SMALL BUT IMPORTANT UX RULE

For photo decisions, the default duration drops to **seven days**, not thirty. Visual choices are easier to feel quickly and lose value if dragged out. Users can override and set longer if they want, but seven is the right anchor.

How the eight formats translate the synthesis differently

The deterministic layers (agreement rate, trajectory) work identically across all eight formats. The AI synthesis prompt is templated to the format: text-based synthesis for Yes/No and Strong/Lean and Scale; image-aware synthesis for Photo A vs B (including the vision descriptions of each image); option-anchored synthesis for Pick Best and Rank Options. The user does not see this. They see one verdict screen that adapts to the format they chose.

SECTION 5

Use cases written like you are twelve

Sometimes the best way to explain a product is to imagine you are explaining it to a twelve-year-old. Every example below uses short sentences, simple words, and concrete situations. No jargon. No business-speak. If a kid would not understand it, the product is not explained well enough.

Couples decisions

Mum and Dad cannot decide if they should move house **REAL BIG DECISION**

Imagine your mum and dad have been talking about moving house for months. Every Saturday they look at houses online. Every Sunday your mum says she would miss the garden. Every Monday your dad says he would love to walk to the cafe. They cannot decide.

So they open Mull Over. Each day for a month, they each tap one button. Strong No, Lean No, Lean Yes, or Strong Yes. Your mum does not see what your dad picked. Your dad does not see what your mum picked. They each also write one sentence about how they feel.

On the last day they look together. Your dad picked Yes nine days in a row. Your mum picked No fifteen times. The app shows them a little paragraph that says, "You are not really split on moving. You are split on what you would give up." Then it asks one question: "Can you find a house in town with a garden?"

Two months later they find one. They move.

Auntie Priya wants a baby. Uncle Tom is not sure yet **FAMILY DECISION**

Auntie Priya is thirty-two and she really wants a baby. Uncle Tom is thirty-one and he is scared about money and not sleeping. They have talked about it many times. They get a bit upset with each other.

They use Mull Over for sixty days, twice as long as normal. Every day they each tap a button. Uncle Tom writes secret notes that he would not say out loud. Things like "I think I am more ready than I told her" and "Maybe I am scared of being scared."

On day sixty they see the verdict. Auntie Priya wants a baby (very strong). Uncle Tom wants one too, but he is only a little bit ready. But here is the surprise: his answers got more yes as the weeks went on. Hers stayed the same.

The app tells them, "Tom is not saying no. He is saying not yet, but he is moving towards yes." It suggests they pick a date six months away and check in again then. They do that. Next year there is a baby.

My cousin Maya wants to quit her job **SOLO DECISION**

My cousin Maya has a really good job in London. She gets paid lots. But some Mondays she comes home and lies on the floor.

She thinks about quitting and working for herself instead. But she is scared. Her friends do not know. Her mum does not know.

So she uses Mull Over by herself. No partner. She just taps every day for thirty days and writes a little note. On bad Mondays she taps Strong Yes (quit). On Fridays when her team has fun, she taps Lean No (stay).

At the end the app shows her: she is mostly leaning yes, but only just. It says "you feel stuck but not unsafe. You are more frustrated by certain projects than the whole job." It asks: "Could you ask for a four-day week and do your own work on Fridays?"

Maya asks her boss. Her boss says yes. She is still there a year later, happier.

More everyday decisions

Which dress should Lily wear to her sister's wedding **PHOTO DECISION**

Lily is going to her sister's wedding next Saturday. She has two dresses she loves. One is green, one is navy blue. She cannot pick.

She opens Mull Over and uploads two photos: dress A and dress B. She invites her best friend Sara to vote too. For seven days, every morning, they each tap which dress they like better today.

At the end of the week, the app shows them: Lily picked navy four times, green three times. Sara picked navy seven times in a row. The note Sara wrote on day six said "the navy will look amazing in the photos and you can wear it again."

Lily wears the navy.

Mr and Mrs Lee cannot agree on what colour to paint the lounge **HOME DECISION**

Mr and Mrs Lee are doing up the house. They cannot agree on what colour to paint the living room. Mr Lee wants a warm cream. Mrs Lee wants a cool grey. They have argued about it three times.

They open Mull Over and upload two photos: one room painted cream, one room painted grey. For ten days, they each tap which one feels more like home today. They write little notes about why.

At the end the app says, "You both like both. Mr Lee picks the grey when he writes about feeling calm. Mrs Lee picks the cream when she writes about cosy evenings. Maybe you can paint different rooms different colours instead of fighting about one."

They paint the lounge grey and the dining room cream. Both are happy.

Jake and Tom want to name their puppy but cannot agree **TINY DECISION**

Jake and Tom just got a puppy. They cannot agree on a name. Jake wants Biscuit. Tom wants Mango. They have a list of six other names too.

They open Mull Over and pick Rank Options. They put all eight names in. For five days, they each drag the names into their favourite order. They cannot see each other's lists.

At the end the app shows them: both of them ranked Biscuit first or second every single day. Mango was Jake's seventh choice. The app says, "You actually agree more than you think. Biscuit was both your top two picks the whole time. Mango was the disagreement."

The puppy is called Biscuit.

Sam wants to break up but she is not sure **HARD DECISION**

Sam has been with her boyfriend for three years. Some days she thinks she should leave. Some days she is happy. She does not know what to do. She does not want to tell her friends.

She uses Mull Over by herself for sixty days. She does not tell her boyfriend. Every day she taps Strong No, Lean No, Lean Yes, or Strong Yes for "should I end this." She writes tiny notes. "Good day. Made dinner together." "Fight about dishes again."

On day sixty she sees: her score is mostly negative but only just. The trend is slightly going more negative. The app says, "The pattern is small recurring frustrations, not a big problem. You are slowly drifting." It suggests: "Before deciding, have one real conversation about the recurring stuff and check in again in a month."

Sam talks to her boyfriend. The conversation is hard. They go to a couples therapist. A year later they are happier than when she started using the app.

Decisions you might not have thought of

Should I get this tattoo **PHOTO DECISION**

Imagine you really want a tattoo but you are scared you will regret it. You have two designs. Open Mull Over, upload both designs as photos, and vote on them every day for two weeks. Invite a friend whose opinion you trust. At the end you see whether you actually like one more, or whether you have been lying to yourself about loving the second one.

Where should we go on holiday **FAMILY DECISION**

Your family is choosing between three places: Bali, Queenstown, or a road trip. Use Mull Over with Rank Options for fourteen days. Every day each family member ranks the three. At the end you find out which one everyone secretly wanted, and which one was being pushed by the loudest voice in the family.

Should I take this job offer **SOLO DECISION**

You have a new job offer but you also like your current job. Use Mull Over solo for fourteen days. Tap a button every day for "should I take it." Write a one-line note about what is making you lean yes or no today. At the end the app will tell you whether your gut is moving towards the new job or away from it.

Should we get a dog **COUPLE DECISION**

You and your partner have been talking about getting a dog for ages. One of you really wants it. One of you is worried about the work. Use Mull Over for sixty days. Long enough that the excitement settles. At the end you find out whether you both want it or just one of you really wants it and is pulling the other along.

Which baby name **PHOTO OR RANK DECISION**

You are six months pregnant. You and your partner have a shortlist of names but cannot agree. Use Mull Over with Rank Options. Put your five favourites in. Each rank them every day for two weeks. The verdict will show you which names were both your top picks, which were one person's pick only, and which were nobody's real favourite even if they were on the list to be polite.

Should we have another baby **COUPLE DECISION**

You already have one or two kids. You cannot agree on whether to have another. Use Mull Over for sixty or even ninety days. This is a decision that needs a long window because how you feel changes with how the kids are sleeping that week. At the end the app shows you the real pattern under the daily ups and downs.

Should we sell this share, this car, this rental property **FINANCIAL DECISION**

You have an asset you are thinking of selling. The price moves around. Your feelings move around. Use Mull Over solo or with your partner for thirty days. At the end you find out whether you genuinely want to sell or whether you are just reacting to the latest news.

Which preschool, which primary school, which high school **PARENTING DECISION**

You are choosing between schools for your child. You have visited two or three. Use Mull Over with Pick Best for fourteen days. Every day, each parent picks which school feels right today. At the end you see whether you both consistently pick the same one (decision made) or whether you keep switching (you need more information).

SECTION 6

Mull Over for Teams: the corporate use case

The same mechanic that helps a couple decide whether to move house also helps a team of six decide which logo to launch with, which feature to build next, or which job candidate to hire. Mull Over for Teams is a B2B-lite extension of the consumer product, sharing the same backend, the same divergence analysis, but with multi-voter support and a small set of additional features tailored to organisational decisions.

The structural similarity

Group decisions in companies suffer from the same flaw as couple decisions: the loudest voice or the most senior voice anchors everyone else. The CEO mentions she likes option B in passing, and by the time the team meeting starts, everyone agrees with option B. The honest signal never reaches the data.

Mull Over for Teams applies the same privacy-locked time-series voting to organisational decisions. A decision-owner (a manager, a designer, a product lead) creates a decision, invites between three and twenty colleagues by email, sets a duration (often seven to fourteen days, sometimes longer), and the team votes daily. Nobody can see anyone else's votes during the period. At the end, the synthesis names the actual pattern: where the team converged, where it diverged, and what the disagreement is really about.

Use cases that fit corporate teams

HIRING DECISIONS BETWEEN TWO FINALISTS

A panel of four interviewers has met both candidates. The traditional approach is a forty-five minute debrief meeting where the most senior interviewer anchors the discussion. Mull Over for Teams alternative: each interviewer votes daily for seven days (during reference checks and reflection period) on a Strong/Lean scale, with notes. At day seven, the panel reveals their actual independent positions. The hiring lead sees patterns: "the candidate everyone said was the right pick had three Lean No votes in week one that softened. The 'maybe' candidate had one Strong Yes that nobody told the room about."

DESIGN AND BRAND DECISIONS

Logo A versus logo B. Brand colour 1 versus brand colour 2. Marketing concept this versus that. Photo format, ten-person team, seven days. At end of period, the synthesis tells the design lead which option had broader appeal versus which had stronger advocates, and where the disagreement clustered (often along function lines: engineers picked one, sales picked another).

PRODUCT STRATEGY DECISIONS

"Which feature should we build next?" with five candidates and a Rank Options format, fourteen days. Engineering, design, sales, and customer success each rank daily. The verdict shows where the rankings converge and where they fracture by function. Common output: "Sales and customer success consistently ranked feature X first. Engineering and design consistently ranked feature Y first. The disagreement is about which kind of value you want to create in the next quarter."

OFFICE AND OPERATIONS DECISIONS

"Should we open a Sydney office?" "Should we switch to remote-first?" "Should we change our health insurance provider?" These are decisions where the group's stated position in a meeting often differs from the group's actual independent opinion. Mull Over for Teams gives the decision-owner real data instead of consensus theatre.

BOARD AND GOVERNANCE DECISIONS

Non-profit boards, co-op governance bodies, and small-company boards often face polarising decisions where the loudest member dominates discussion. Anonymous longitudinal voting before the board meeting surfaces the actual distribution of opinion and prevents groupthink. This use case has direct overlap with Loomio, which is the closest current alternative (see section 10).

What is different from the consumer product

MULTI-VOTER, NOT TWO-VOTER

Consumer Mull Over caps at two participants per decision. Teams supports between three and twenty voters per decision. The synthesis adapts: instead of "you are not split on X, you are split on Y" for two partners, it reads "the team converged on X, with the disagreement clustered around Y. The strongest dissent came from a sub-group whose notes mentioned [theme]."

PER-TEAM SUBSCRIPTION, NOT PER-DECISION

Per-decision pricing works for couples (one decision a year is common, paying twenty dollars for a meaningful answer feels right). Companies make decisions monthly or weekly, and the procurement process for one-off twenty-dollar charges is friction. Teams pricing should be per-seat or per-team monthly subscriptions.

DECISION OWNER ROLE, NOT EQUAL PARTICIPANTS

Couples are equal participants. Team decisions have an owner (the manager, the designer, the hiring lead) who frames the question, invites the voters, sees the verdict, and decides what to do with it. The owner has a small amount of additional visibility (e.g., overall participation rates per voter) but never sees individual votes during the period.

INTEGRATION WITH TEAM TOOLS

Slack notifications (slash command to create a decision, daily DMs to vote, verdict posted to channel). Microsoft Teams integration for enterprise. Single sign-on with Google Workspace and Okta for organisations larger than ten seats. These are the standard B2B trappings, important for adoption inside companies.

ANONYMISATION AT THE VERDICT, WITH OPT-OUT

By default, the verdict screen shows team-level patterns, not individual votes. ("Eight team members leaned yes, three leaned no, one strong no.") The synthesis describes themes from notes without attributing them. Decision owners can opt into named-vote verdicts for use cases where attribution matters (formal board votes, governance), but the default is anonymous.

Pricing for teams

TIER	PRICE	INCLUDES
Teams Starter	\$49 per month	Up to 10 voters, unlimited decisions, all eight formats, Slack integration, divergence analysis on every decision
Teams Growth	\$149 per month	Up to 50 voters, all of above, MS Teams integration, decision templates library, named-vote opt-in
Teams Enterprise	\$499 per month	Unlimited voters, SSO, audit log, dedicated support, custom synthesis prompts, on-call for high-stakes decisions

Why this matters for the business case

The consumer side of Mull Over is a small business with strong unit economics but capped scale (a couple does not make many decisions a year). The team side has different shape: smaller customer count, much higher revenue per customer, recurring rather than one-off. Five hundred Teams Starter customers at forty-nine per month is two hundred and ninety-four thousand dollars in annual recurring revenue, which by itself meets the "one hundred thousand per year" threshold this document analyses in section 23.

The recommended sequence is to ship the consumer product first (the couples version validates the synthesis quality and brand), then layer the Teams version in month nine or ten once the AI synthesis is proven. The technical lift to add multi-voter and Slack integration on top of the existing backend is small: roughly four to six weeks of engineering.

A NOTE ON POSITIONING

"Mull Over for Teams" is the working name but is not necessarily the launch name. The brand may benefit from a separate B2B identity (e.g., *Converge*, *Verdict*, or simply *Mull Over Team Edition*) to avoid B2B buyers being put off by the couples app association. This is a marketing decision to be tested closely on launch, not a product decision.

The market is real. The competition is broader

SECTION 7

competition is broader Market sizing and the realistic addressable slice than it looks.

The couples app market is large and growing. The realistic addressable slice for Mull Over inside that market is small and high margin. This section gives a full summary of the market honestly rather than the top-down market sizing that often looks impressive in fundraising decks but is unhelpful for an operating founder. Technology, couples therapy, decision-making tools, and photo voting. Plus the feature matrix and a structural explanation of why the gap has stayed open.

\$2.0B

COUPLES APPS IN 2024

\$5.77B

FORECAST 2033

12.5%

CAGR 2025 TO 2033

Source: Business Research Insights, 2025. This is the broad couples app category, including communication tools, shared-space utilities, daily-question apps, intimacy apps, and counseling apps. Mull Over fits inside this category but does not compete head-on with most of it.

The couples counseling app sub-segment

\$720M

2025 SIZE

\$2.1B

2032 FORECAST

14.2%

CAGR 2025 TO 2032

Source: HTF Market Intelligence, March 2026. The faster growth in this sub-segment reflects two trends: increasing acceptance of digital relationship support after the pandemic, and the rising cost and access difficulty of in-person couples therapy.

The realistic addressable slice for Mull Over

Top-down market sizing is unhelpful for the founder's purposes. The more honest exercise is bottom-up. Mull Over does not need to capture even a fraction of a percent of the couples app market to be a successful business. What it needs is a few thousand paying customers per year in the first three years.

The five Tier 1 English-language markets (US, UK, Canada, Australia, New Zealand) contain roughly one hundred million couples in the prime usage demographic of twenty-five to fifty-five years old. The bottom-up funnel:

- **Awareness:** 0.5% to 2% of target couples hear about Mull Over within twenty-four months of launch (depending on marketing intensity)
- **Trial:** 5% to 10% of those aware sign up for a free solo decision

- **Shared decision:** 12% to 18% of solo triallers progress to a shared decision within ninety days
- **Per-decision purchase:** 55% to 65% of shared decisions complete and pay for divergence analysis
- **Annual upgrade:** 20% to 30% of per-decision purchasers upgrade to annual within twelve months

Plugging the middle of these ranges into the US market alone (seventy million target couples) gives twenty-four-month numbers of approximately three thousand five hundred paying customers, or roughly six hundred thousand dollars in cumulative revenue across the period. Blended across all five Tier 1 markets, twenty-four-month customer count lands in the range of two thousand to ten thousand, with annual run-rate revenue between sixty thousand and four hundred thousand dollars depending on which assumptions inside each range play out.

These numbers are deliberately more conservative than the first edition of this document. Edition one suggested ranges up to two point one million dollars ARR at month twenty-four. That was a real possibility under aggressive assumptions but it required multiple major channels to land simultaneously. Edition two anchors on the more probable middle-of-the-distribution outcomes.

Why the realistic slice still works

Even at the low end of the funnel assumptions, the business is viable because the cost base is so small. Section 14 details the pricing. Section 15 details the unit economics. Section 22 details the three revenue scenarios with full P&L. The point of this section is to anchor expectations: this is a small-but-profitable opportunity, not a venture-scale one. The founder should build it on that basis.

Direct competitors in couples technology

Seventeen named products serve the couples segment of the relationship app market. None offer Mull Over's specific combination of features, but each one fights for attention from the same prospective user. Knowing each competitor in detail shapes positioning and marketing.

Paired

\$8.99/mo, \$71.99/year

LONDON-BASED · 8 MILLION DOWNLOADS · LAUNCHED 2020

The leading daily-question app for couples. Each partner gets a daily prompt, answers, then sees the other's answer. Topics include intimacy, communication, finances, family. Marketed as relationship maintenance, not therapy. Backed by Notion Capital and angel investors.

How Mull Over differs: different job to be done. Paired is daily relationship maintenance. Mull Over is specific decision support. A couple might use both.

Lasting

\$11.99/mo, \$119.99/year

ACQUIRED BY THE KNOT · 3+ MILLION COUPLES · GOTTMAN-BASED CURRICULUM

Structured marriage counseling app built around the Gottman method. Five-minute daily exercises across conflict, intimacy, finance, family. Acquired by The Knot in 2020 and tightly integrated with the wedding-planning funnel. 94% of users report stronger relationships.

How Mull Over differs: Lasting is a therapy curriculum. Mull Over is decision support. Lasting teaches you communication skills; Mull Over gives you data about a specific decision.

Relish

\$19.99-\$29.99/mo, \$240-\$360/year

HELLORELISH LTD · COACHING + CONTENT

Relationship coaching app with personalised plans and optional access to human coaches via chat. Psychology-based quizzes that decode each partner's thinking. Premium tier with weekly coach sessions positioned as therapy-adjacent.

How Mull Over differs: Relish is human-in-the-loop coaching. Mull Over is software-only with AI synthesis substituting for the coach. Different price tier, different mechanic.

Coral

\$7.99/mo, \$95.99/year

SEXUAL WELLNESS FOCUS · STRONG CONTENT LIBRARY

Focused specifically on intimate and sexual wellness for couples. Less competition in this sub-niche than in general couples communication. Content-heavy product, low daily-active expectation.

How Mull Over differs: different domain entirely. Coral is intimacy-focused content; Mull Over is decision support. No overlap in core use case.

OurRitual

\$32+/week, \$1,664+/year

HYBRID COACH + STRUCTURED PATHWAY · PREMIUM TIER

Weekly twenty-to-forty-minute virtual sessions with a trained relationship expert, plus a library of guided modules ("pathways") that the couple works through between sessions. Premium-tier alternative to traditional therapy.

How Mull Over differs: OurRitual is premium guided therapy. Mull Over is the affordable, decision-specific complement at one-twentieth the price.

Between

Free with in-app purchases

KOREAN, BY VCNC · 35 MILLION COUPLES GLOBALLY

Private messaging and shared-space app for couples. Photos, messages, anniversaries, shared calendar. The largest couples app globally by user count, particularly dominant in Korea and Japan.

How Mull Over differs: Between is a utility (private space for two). No decision logic, no AI insight, no structured voting. Different category entirely.

OurCouple

Free with premium tiers

IOS-FIRST · COMBINED JOURNAL AND SHARED SPACE

Newer entrant combining daily messages and photos, shared memory journal, love notes, shared event calendar, and a virtual pet shared with the partner. Positions itself as the all-in-one private space for two.

How Mull Over differs: OurCouple is breadth (does many things). Mull Over is depth (does one thing well). No overlap on decision support.

Connected

Free with paid tiers around \$9-12/mo

ALL-IN-ONE RELATIONSHIP PLATFORM

Positioned as overall-relationship-health platform: daily check-ins, shared goals, relationship insights, communication exercises. Strong free tier with paid upgrades for analytics and coaching content.

How Mull Over differs: Connected is broad relationship maintenance. Mull Over is narrow decision support. The two would coexist if both were on a couple's phone.

Lovewick

\$9.99/mo, \$59.99/year

DATE NIGHT IDEAS, ROMANCE, APPRECIATION

Focused on keeping romance alive: date-night suggestions, appreciation prompts, anniversary tracking. Lighter-touch than therapy-adjacent products.

How Mull Over differs: different problem. Lovewick is romance. Mull Over is decision-making.

Agapé

\$59.99/year

DAILY QUESTIONS · RESEARCH-BACKED

Built around twenty years of relationship research. Daily questions for emotional intimacy and curiosity-building. Often recommended alongside Gottman Card Decks for couples doing structured emotional work.

How Mull Over differs: daily-question maintenance product. No decision logic.

Love Nudge

Free

BASED ON GARY CHAPMAN'S 5 LOVE LANGUAGES

Gamified love-language tracking. Each partner identifies their love language and the app nudges the other to do small things in that language. Strong organic discovery via the 5 Love Languages book franchise.

How Mull Over differs: different framework, different problem. Love Nudge is about giving; Mull Over is about deciding.

Coupleness

\$4.99/mo, \$39.99/year

RELATIONSHIP FITBIT · DAILY MOOD LOGGING

Daily relationship mood logger. Each partner taps how they feel about the relationship today. Visual graphs of trends over time. Lightweight, low-friction, low-stakes.

How Mull Over differs: closest in mechanic (daily tap, time-series). But Coupleness is open-ended (no specific decision), no AI synthesis, partner sees the data immediately. The privacy mechanic and decision-anchoring is the differentiator.

Mend

\$2.99/mo, \$14.99/year

POST-BREAKUP RECOVERY, INDIVIDUAL USE

Heartbreak coach for individuals after breakups. Daily exercises, audio content, journaling. Solo product, not for couples.

How Mull Over differs: different stage of relationship. Mend is post-relationship. Mull Over might be used to decide whether to end one (Sam's story, page 14), but the use case is different.

Cupla

\$3.99/mo

COUPLES CALENDAR · IOS ONLY

Shared calendar for two. Syncs with Google and Apple calendars and shows both partners' schedules side by side. Single-purpose utility for coordinating schedules without arguing.

How Mull Over differs: utility, not decision support.

Raft

Free with premium

CALENDAR APP FOR COUPLES AND CLOSE FRIENDS

iOS calendar app for couples and close groups. Planning, day-coordination, shared events. Single-purpose.

How Mull Over differs: utility, not decision support.

Honeydue

Free with premium

COUPLES FINANCE APP

Joint budgeting, bill tracking, and financial conversation prompts for couples. Acquired and integrated into broader fintech offerings.

How Mull Over differs: finance utility, not decision support. Honeydue might appear adjacent to a "should we buy this house" decision but does not provide the decision logic.

Donkey Chats

\$49/year

AI CONFLICT ANALYSIS · RECORDS ARGUMENTS

Newer AI-powered relationship app: records arguments and analyses tone, interruptions, emotional spikes, and speaking balance. Positioned as conflict de-escalation tool.

How Mull Over differs: Donkey Chats analyses conflict in real time after it happens. Mull Over captures positions before the conflict. Different point in the decision lifecycle.

LoveFix

Free trial, \$9.99/mo

IN-THE-MOMENT CONFLICT SUPPORT

Real-time guidance during difficult conversations. Designed to be used mid-argument to suggest scripts and de-escalation techniques. Newer entrant with strong content marketing presence.

How Mull Over differs: in-the-moment conflict, versus pre-decision structured capture. Complementary, not competitive.

Summary of the couples-technology landscape

The pattern is clear: seventeen products competing for the same prospective couple, but each occupying a different sub-niche. Daily-question apps (Paired, Lasting, Agapé). Therapy curricula (Lasting, OurRelationship). Shared spaces and utilities (Between, OurCouple, Cupla, Raft, Honeydue, Coupleness). Coaching (Relish, OurRitual). Specific domains (Coral for intimacy, Mend for post-breakup, Love Nudge for love languages). AI conflict analysis (Donkey Chats, LoveFix).

None occupy the decision-aggregation slot. The closest in mechanic (daily tap, time-series) is Coupleness, but it is open-ended mood tracking without the structural ingredients that make Mull Over work: a specific decision, a fixed duration, privacy until verdict, AI synthesis. The closest in price tier is Lasting at one hundred and twenty per year, but Lasting is therapy curriculum, not decision support.

The takeaway for positioning: Mull Over should not market itself as "another couples app." That framing puts it in a saturated category and forces direct comparison against Paired or Lasting on dimensions where neither is the right match. Mull Over should market itself as a "decision aggregator," which is its own category, distinct from the existing seventeen.

SECTION 9

Couples therapy apps and clinical platforms

A separate tier of competitors operates as clinical therapy delivery: licensed therapists, weekly video or messaging sessions, prescription-grade pricing. These are not direct competitors for Mull Over but they shape what couples consider when they have a problem, and the Mull Over founder needs to understand this market in detail.

BetterHelp Couples

\$240-\$400/mo, \$2,880-\$4,800/year

LARGEST ONLINE THERAPY PLATFORM · LICENSED THERAPISTS

The dominant online couples therapy provider. Match with a licensed couples therapist, unlimited messaging, optional weekly video sessions. Backed by Teladoc Health. Aggressive marketing budget; ubiquitous in podcasts and YouTube ads.

How Mull Over differs: therapy is human-delivered ongoing conversation. Mull Over is software that prepares the conversation. Many therapists would recommend Mull Over to clients between sessions.

Talkspace

\$276-\$396/mo

PUBLIC COMPANY · INSURANCE-FRIENDLY

Online therapy platform with strong insurance integration. Couples therapy add-on layered onto individual therapy. Different positioning from BetterHelp (more clinical, less mass-market).

How Mull Over differs: same as BetterHelp.

ReGain

\$240+/mo

BETTERHELP SUBSIDIARY, COUPLES-FOCUSED

BetterHelp's dedicated couples therapy platform. Same licensed-therapist model with couples specialisation. Often appears in Google search when couples search for online therapy.

How Mull Over differs: same.

OurRelationship

Free with paid premium

RESEARCH-BACKED · UCLA / BROWN CLINICAL ROOTS

Evidence-based structured program developed by university researchers. The strongest direct clinical trial evidence in the couples-app space. Substantial free tier; the paid premium adds expert support.

How Mull Over differs: structured therapy course versus decision support. Complementary; OurRelationship users could happily use Mull Over for specific decisions.

Online-Therapy.com

\$48-\$96/week, \$2,500-\$5,000/year

CBT-BASED · WORKSHEET-HEAVY

Cognitive Behavioural Therapy-focused platform with worksheets, journal prompts, and thought logs. Couples plan starts at ninety-six dollars per week.

How Mull Over differs: CBT-driven therapy versus structured decision support.

Calmertry

\$280-\$360/mo

NEWER ONLINE THERAPY PLATFORM

Modern online therapy platform with couples option. Smaller than BetterHelp but growing fast in 2025-2026.

How Mull Over differs: therapy delivery, not decision support.

Thriveworks

\$160-\$300/session

INSURANCE-FRIENDLY · IN-PERSON AND ONLINE

Large clinical network in the United States with both in-person and online therapy. Insurance accepted. Couples therapy is a major service line.

How Mull Over differs: clinical therapy versus software.

Growing Self

\$120-\$200/session

ONLINE AND IN-PERSON THERAPY

Specialised couples and family therapy provider. Strong content marketing presence with relationship advice blog. National US coverage.

How Mull Over differs: clinical therapy versus software.

Couples Therapy Inc.

\$200-\$400/session, intensive weekends \$4,000-\$8,000

INTENSIVE WORKSHOPS + ONGOING THERAPY

Premium couples therapy practice known for intensive weekend retreats. Targets couples in crisis. Higher price point, lower volume.

How Mull Over differs: crisis-tier intervention versus pre-decision support.

Amwell, MDLive, Teladoc

\$80-\$300/session

TELEHEALTH PLATFORMS WITH THERAPY OPTIONS

Large telehealth platforms that include therapy as one service among many. Couples therapy is available but not the focus. Insurance-friendly.

How Mull Over differs: general telehealth versus relationship-specific decision software.

Why the therapy tier matters for Mull Over positioning

The therapy tier is fifty to one hundred times more expensive than Mull Over and serves a different need (ongoing conversation with a trained professional). But the therapy tier is also Mull Over's most natural channel partner. Couples therapists face a recurring problem: clients arrive without structured data about their actual disagreement, so the first three sessions are spent extracting that data manually. A therapist who can hand a client Mull Over before the first session shortens their work and produces better outcomes for the client. This is the basis of the therapist partnership program detailed in section 20.

Pricing context

TIER	ANNUAL COST	WHAT YOU GET
Mull Over (annual)	\$99	Unlimited shared decisions with divergence analysis
Paired (premium)	\$72	Daily questions, expert content
Coral (premium)	\$96	Intimacy-focused content
Lasting (premium)	\$120	Gottman therapy curriculum
Relish (with coach)	\$240-\$360	Coaching plus content
OurRitual	\$1,664+	Weekly sessions with relationship expert
BetterHelp Couples	\$2,880-\$4,800	Licensed therapist, weekly video
Traditional couples therapy	\$6,000-\$15,000	In-person therapy, weekly

Mull Over sits at the entry point of the price ladder, deliberately. It is software with one specific job, and that job costs roughly one decision out for dinner per shared decision.

SECTION 10

Decision-making and group polling tools

A distinct competitor category exists for group decision-making, mostly enterprise or community-focused rather than consumer. These products are not direct competitors to consumer Mull Over but they are very direct competitors to Mull Over for Teams (section 6). Knowing them in detail shapes the B2B positioning.

Loomio

Free for small groups, \$25-\$300/mo for organisations

OPEN-SOURCE · NZ-FOUNDED · USED BY COOPERATIVES, NON-PROFITS, UNIONS

The closest existing tool to Mull Over for Teams. Open-source collaborative decision-making platform. Supports multiple poll types (proposal, multiple choice, ranked choice, score, dot voting, time poll). Has anonymous voting and "hidden results" mode where results are not visible until the poll closes (this is the closest mechanic to Mull Over's privacy lock). Used widely by community groups, non-profits, and small businesses.

How Mull Over Teams differs: Loomio is a single-event poll tool (all votes happen at one point in time before a deadline). Mull Over Teams is longitudinal (votes happen daily over a period). The synthesis layer is the bigger differentiator: Loomio shows you who voted what; Mull Over Teams tells you what the disagreement is actually about.

Cloverpop

\$300+ per user per month, enterprise

ENTERPRISE DECISION SOFTWARE · B2B ONLY

Enterprise decision-making platform for executive teams. Requires at least four options per decision and at least five participants. Includes decision tracking, audit trail, and decision-quality scoring. Used by large corporates for executive-level decisions.

How Mull Over Teams differs: Cloverpop is enterprise-priced (three hundred dollars per user per month versus forty-nine dollars per month for a team of ten). Cloverpop also lacks longitudinal capture and AI synthesis.

Polly

Free, \$19-\$39/mo

SLACK-NATIVE POLLING

The dominant Slack-integrated polling tool. Create polls directly in Slack channels, vote inline. Templates for retros, retrospectives, sentiment surveys. Strong for instant team polls but lacks decision-specific structure or longitudinal capture.

How Mull Over Teams differs: Polly is instant single-event polls. Mull Over is a thirty-day deliberate process. Different problem.

Doodle

Free, \$6.95-\$30/mo

TIME AND DATE POLLING

The classic group polling tool, primarily for scheduling meetings. Can be used for general polling but the product is heavily oriented toward time-and-date coordination. Used by tens of millions globally.

How Mull Over Teams differs: Doodle is a scheduling utility, not a decision-making tool.

1000minds

\$50-\$500/mo

CONJOINT ANALYSIS · NZ-FOUNDED

Sophisticated multi-criteria decision-making software using conjoint analysis. Surfaces relative preferences across many factors. Used heavily in healthcare and policy for prioritisation. Powerful but with a steep learning curve.

How Mull Over Teams differs: 1000minds is for sophisticated multi-attribute decisions where you need to weight many factors. Mull Over is for the human-level question of "where are people actually at on this." Different sophistication tier.

Decisive

Free, \$3.99 unlock

CONSUMER MOBILE APP · SINGLE-USER

Consumer-facing decision-making app for individuals. Pros and cons format with weighted scoring. Single-user, single-decision. Small footprint, low engagement.

How Mull Over differs: Decisive is one-shot solo pros-and-cons. Mull Over is multi-day, multi-person, with AI synthesis. Different category of product.

ChoiceMap

Free, \$9.99 premium

DECISION-MAKING FRAMEWORK APP

Consumer-facing decision tool implementing the analytic hierarchy process. Helps weigh multiple options across multiple criteria. Solo use, single-session.

How Mull Over differs: framework-based single-session versus longitudinal multi-person.

Decision Journal

\$6.99 one-off

SOLO DECISION JOURNAL · BEZOS-INSPIRED

Inspired by Jeff Bezos's annual letters about decision-making, this app captures your decisions, expected outcomes, and confidence levels. Lets you review past decisions to improve over time. Solo use only.

How Mull Over differs: retrospective decision journaling versus prospective decision support.

Mentimeter, Slido, Kahoot

\$10-\$300/mo

INTERACTIVE PRESENTATION POLLING

Real-time audience polling tools for presentations and meetings. Strong in education and corporate event use cases. Not designed for asynchronous longitudinal decision-making.

How Mull Over differs: live-event polling versus thirty-day deliberation.

What this category tells us about Mull Over for Teams positioning

Three observations:

First, every existing decision-making tool is either point-in-time (Doodle, Polly, Loomio, Mentimeter) or solo-use (Decisive, ChoiceMap, Decision Journal). Longitudinal multi-participant capture is genuinely uncontested. The closest is Loomio's "hidden results" mode, but Loomio is still a single-event poll where everyone votes once.

Second, the pricing distribution shows a clear gap. Consumer tools are five to ten dollars. Polly and Loomio are in the twenty to fifty dollar per month range. Cloverpop is three hundred plus per user per month. Mull Over Teams at forty-nine, one hundred and forty-nine, and four hundred and ninety-nine dollars per month per team (not per user) fits a small-business and mid-market gap that is currently underserved.

Third, the AI synthesis layer is the most defensible single feature. No competitor in this category has it. Loomio comes closest with their "synthesis" feature, but it is structured discussion summarisation, not pattern extraction from longitudinal voting data. Loomio's value proposition is "we make the discussion organised"; Mull Over Teams' is "we tell you what the data actually says, which is often different from what the discussion concluded."

SECTION 11

Photo voting and image polling tools

A fourth competitor category exists for visual A-versus-B decisions. These products mostly target either e-commerce (testing product photos) or casual social polling. None offer the longitudinal-voting plus AI-synthesis combination that Mull Over brings, but they are the closest tools couples currently use when they need to decide between two outfits or two paint colours.

PollUnit

Free with paid tiers \$12-\$45/mo

IMAGE POLLS AND PAIRWISE COMPARISON

The most comprehensive image polling tool. Supports plain image polls, ranked image polls, and pairwise comparison (good for evaluating many images by comparing two at a time). Used by photographers, designers, and small businesses for customer feedback.

How Mull Over differs: PollUnit is single-event, public-or-invited polls. No longitudinal capture, no AI synthesis, no privacy mechanic during voting. Suited to crowd opinion; not suited to private personal decisions.

Opinion Stage

Free, \$25-\$220/mo

QUIZ AND POLL BUILDER

Comprehensive quiz and poll builder, often embedded in marketing landing pages. Supports image-based polling alongside text quizzes. Used heavily by content marketers for engagement.

How Mull Over differs: marketing engagement tool, not decision support.

Vote Social

Free

CONSUMER MOBILE APP FOR A-VS-B PHOTO POLLS

Mobile app where users upload two photos and crowdsource opinions from friends, followers, or the public Vote Social community. Like Instagram polls but standalone. Strong in style decisions, "rate my outfit" use cases.

How Mull Over differs: public crowdsourcing versus private longitudinal voting. Vote Social is "should I post this photo"; Mull Over Photo is "which of these two should I commit to."

Woobox

Free with paid tiers \$37-\$249/mo

PHOTO CONTEST AND VOTING PLATFORM

Marketing-oriented photo contest builder. Brands run "submit your photo, public voting picks the winner" campaigns. Used heavily in CPG marketing.

How Mull Over differs: brand marketing contest tool, not decision support for individuals or couples.

Traxoft (Shopify app)

\$9.99-\$49.99/mo

E-COMMERCE A-B TESTING FOR PRODUCT IMAGES

Shopify app for online stores to validate product images and ideas through customer voting. Strictly B2B SaaS for e-commerce stores.

How Mull Over differs: different industry, different buyer.

Instagram Polls and TikTok Polls

Free, built into platforms

STORIES-BASED POLLING FOR FOLLOWERS

The free default that most casual photo polls happen on today. Post two photos to your Story, ask followers to pick, results disappear in twenty-four hours. Hugely used but limited: public to followers, no longitudinal capture, no analysis.

How Mull Over differs: Instagram polls are public-to-followers single-event. Mull Over Photo is private-to-couple longitudinal. The privacy and the day-by-day capture matter for higher-stakes decisions.

What this means for the photo format

Photo voting is a real existing behaviour. Couples already do it on Instagram, on group chats, on Vote Social. The opportunity is not to invent photo voting; it is to give the existing behaviour the structure that makes the outcome more useful: privacy until the end, daily capture rather than single moment, AI synthesis that explains what each person was actually weighing. The format expansion in section 4 is a modest engineering addition that opens a category of casual everyday decisions Mull Over would otherwise miss.

Feature matrix and where Mull Over wins

Across the four competitor categories from sections 8 to 11, one consistent pattern: nobody offers Mull Over's specific combination of features. The matrix below makes this concrete. It is not exhaustive (it would be unreadable at thirty products); it picks one representative from each category plus the two closest mechanics.

FEATURE	MULL OVER	PAIRED	LASTING	BETTERHELP	LOOMIO	COUPLENESS	POLLUNIT
Decision-specific framing	✓	—	—	~	✓	—	✓
Two-person couple mode	✓	✓	✓	✓	~	✓	—
Multi-person team mode	✓	—	—	—	✓	—	✓
Privacy-locked partner data	✓	—	—	—	✓	—	—
30-day longitudinal capture	✓	—	—	—	—	✓	—
AI divergence synthesis	✓	—	—	—	—	—	—
Themes from notes (frequency)	✓	—	—	—	—	—	—
Conversation prompt at verdict	✓	—	—	~	—	—	—
Photo voting format	✓	—	—	—	—	—	✓
Per-decision pricing option	✓	—	—	—	—	—	—
Branded PDF verdict export	✓	—	—	—	✓	—	~
Daily relationship questions	—	✓	✓	—	—	—	—
Therapy curriculum	—	—	✓	✓	—	—	—
Licensed human therapist	—	—	—	✓	—	—	—
Annual price (USD)	\$99	\$72	\$120	\$2,880+	\$300+	\$40	\$144+

Six rows are unique-to-Mull-Over green checks: privacy-locked partner data, thirty-day longitudinal capture, AI divergence synthesis, themes from notes, conversation prompt at verdict, per-decision pricing option. Three of those (privacy, synthesis, per-decision pricing) are the core structural ingredients that nobody else has any natural reason to build, for the reasons section 13 explains.

Why the gap is real and has stayed open

A reasonable reader at this point might ask: if the opportunity is so visible, why has no incumbent built this product? Four structural reasons, each of which alone would keep an incumbent away, and which together keep the gap genuinely open for at least the next eighteen to twenty-four months.

Reason one: the privacy mechanic actively breaks engagement metrics

Every couples app on the market is optimised for daily active use. Their funding, their KPIs, their roadmaps push toward making partners interact more often inside the app. A feature that says "you cannot see your partner's input until day thirty" is heretical to that thesis. Paired's product team would instinctively soften it ("see each other's votes after twenty-four hours, but with a notification") and in doing so destroy the mechanic that makes the data valuable. This is not a hypothetical risk. The privacy mechanic was the single feature most challenged by every reviewer of the early product spec. The answer requires careful explanation. Incumbent product teams do not have time for that explanation.

Reason two: the use case is too narrow for generalists

"Help couples decide on big questions" is a specific use case. Most couples have one or two big decisions per year, not one a week. From the perspective of a generalist couples app aiming for daily active use, this is unattractive. From the perspective of Mull Over, it is perfect: low session frequency but high willingness to pay per session.

The same logic that makes Calm valuable as a sleep app (used at night, not all day) makes Mull Over valuable as a decision app (used during a decision, not all year). Generalist apps struggle to build "moments of high-value, low-frequency use" because their entire engagement model assumes daily attention.

Reason three: the build is non-trivial despite looking simple

Surface-level, Mull Over looks like "an app that asks two people a question every day for thirty days then summarises." That description undersells the technical depth: the partner-invitation flow with magic links and session sync, the privacy-locked database design that genuinely prevents one user from seeing the other's data (not just hiding it in the UI), the deterministic verdict computation across eight question formats, the integration of Claude with cached prompts and structured output parsing, the PDF generation, the multi-platform push notifications, and the trust-building UI work to convince users the votes are actually sealed.

An incumbent with a small team could build this in four to six months. A solo founder with Claude Code support builds it in ten to fourteen weeks. The math favours the founder: by the time an incumbent decides this is worth building, Mull Over is already six or more months ahead in users and brand.

Reason four: the AI synthesis is the moat, and the moat compounds

The deterministic layers of the verdict are commodity. Anyone can compute agreement percentages and regression slopes. The synthesis paragraph and conversation prompt are not commodity. They depend on prompt quality, which depends on iteration against real cases, which depends on having real users to iterate against. The first one hundred paying decisions are reviewed manually and the prompt is tuned. The next nine hundred are evaluated against rubrics derived from the first hundred. By decision one thousand, the synthesis quality is meaningfully harder to replicate than at decision ten.

An incumbent launching a "decision mode" feature in month eighteen does so without the iteration data, without the rubric, without the brand association ("Mull Over is the app that nails the synthesis"). They could catch up, but only by burning their existing user trust on a feature that fails its early tests.

THE WINDOW

Mull Over has roughly an eight-month head start before any incumbent could credibly ship a competing feature. The right way to use that window is not to grow fast at any cost. It is to nail the synthesis quality and build the brand association between "decision aggregator" and "Mull Over" so tightly that even a competing feature from a bigger player feels derivative.

The business: pricing, unit

economics, and where to launch.

Pricing model for consumer and teams

Two parallel pricing structures: a three-tier model for consumers (couples and solo users) and a three-tier model for teams. Each is calibrated against established benchmarks in its segment.

Consumer pricing

TIER	PRICE	WHAT YOU GET
Free	\$0	One active solo decision at a time, all eight vote formats, up to 90 days, deterministic verdict (no AI synthesis), CSV export, PWA install. Generous on purpose.
Per-decision	\$19.99 USD	Charged at verdict, only if the user wants to unlock divergence analysis. The 30 days of voting are free. Adds AI synthesis, themes, conversation prompt, branded PDF.
Annual	\$99 USD/year	Unlimited shared decisions, up to 5 in parallel, durations up to 365 days, priority push, priority support, credit for past per-decision purchases.

Two design choices worth flagging. First, the free tier is generous because most users who try Mull Over solo are evaluating it; they need to feel the verdict quality on a stakes-low decision before they share with a partner. Second, per-decision pricing is charged *at verdict*, not upfront. The thirty days of voting cost zero. Users only pay when the AI insight is ready. This is unusual but intentional: it lets people complete the full mechanic before deciding whether the synthesis is worth twenty dollars.

Teams pricing

TIER	PRICE	VOTERS	NOTABLE FEATURES
Teams Starter	\$49/mo	Up to 10	Unlimited decisions, all 8 formats, Slack, divergence analysis on every decision
Teams Growth	\$149/mo	Up to 50	Plus MS Teams, decision templates library, named-vote opt-in, prioritised support
Teams Enterprise	\$499/mo	Unlimited	SSO (Okta, Google Workspace), audit log, custom synthesis prompts, dedicated on-call

Teams pricing is per-team-per-month, not per-seat. Forty-nine dollars per month is below the threshold where most teams need to engage procurement. Lower-tier teams can simply expense it. The four-hundred-and-ninety-nine dollar tier is where Mull Over starts to compete more directly with Cloverpop (three hundred dollars per user per month for enterprise decision software), at a meaningful discount.

Why these specific numbers

WHY \$19.99 FOR PER-DECISION

Three reference points anchor this number. First, a Headspace single-course pack at nineteen ninety-nine; consumers know what this price feels like and what to expect from it. Second, the willingness-to-pay anchor for a single specific stakes-high decision: people will pay twenty dollars to get a clear answer on a question they have been wrestling with for months, but resist a thirty dollar charge. Third, the Stripe per-transaction fixed fee (thirty cents) eats four percent of the gross at twenty dollars, fifteen percent at five dollars; nineteen ninety-nine puts the math in margin-friendly territory.

WHY \$99 FOR ANNUAL

Sits cleanly between Paired (seventy-two dollars) and Lasting (one hundred and twenty dollars). Frames as "the price of five per-decision purchases," which is the breakeven point for the annual to make sense. Below the psychological one-hundred-dollar threshold but signals premium against the seventy-two-to-eighty dollar mass-market tier.

WHY \$49 FOR TEAMS STARTER

Below the procurement threshold for most companies. Below the SaaS subscription threshold where finance reviews kick in. Comparable to Polly's mid-tier or a small Loomio plan. The team of ten that adopts Mull Over does it because one team member tried it and pulled the others in, not because procurement approved it.

Unit economics, decomposed

The unit economics are the foundation of the business case. Margins are intentionally high because the cost structure is minimal: no human in the loop, no inventory, no logistics, and AI inference cheap enough to be effectively free per transaction.

Per shared decision (\$19.99 charge)

Gross revenue	\$19.99
Stripe fee (US domestic, 2.9% + \$0.30)	\$0.88
Claude API (Sonnet 4.6, with caching)	\$0.03
Email delivery (Postmark, ~6 emails per decision)	\$0.006
PDF generation (compute on Hetzner)	\$0.001
Net contribution per decision	\$19.07
Effective margin	95.4%

Per annual subscriber (\$99 charge)

Gross revenue	\$99.00
Stripe fee (US, with Billing add-on 0.7%)	\$3.86
Claude API (avg 4 shared decisions/year)	\$0.12
Email delivery (~30 emails/year)	\$0.03
Compute and storage (amortised)	\$0.50
Net contribution per annual sub	\$94.49
Effective margin	95.4%

Per Teams Starter subscription (\$49/mo, \$588/year)

Gross revenue (annual)	\$588.00
Stripe fees (12 monthly charges + Billing add-on)	\$24.59
Claude API (avg 24 team decisions/year)	\$0.72
Email and push delivery	\$2.40
Compute and storage	\$3.00
Slack and MS Teams API costs	\$0.00
Net contribution per Teams Starter	\$557.29
Effective margin	94.8%

Lifetime value modelling

LTV varies by acquisition cohort. Three plausible consumer cohorts:

- **Cohort A: Per-decision only.** Buys one decision, never returns. LTV \$19.07.
- **Cohort B: Per-decision then annual upgrade.** Buys one decision (contribution \$19.07), upgrades to annual after sixty days with credit applied (\$74.94 net), retains for an average of 1.6 years. Blended LTV approximately \$145.
- **Cohort C: Direct annual.** Skips per-decision. Retains 1.8 years. LTV approximately \$170.

Weighting the consumer cohorts (estimated 45% / 35% / 20%) gives a blended consumer LTV of approximately **\$89 per paying user**.

For teams, LTV is meaningfully higher per customer because the subscription is recurring. Teams Starter retained for an average of 14 months gives an LTV of \$650. Teams Growth retained 16 months gives \$2,384. Teams Enterprise retained 24 months gives \$11,976. A small number of Teams customers can move the revenue line considerably (this is the basis for the probability table in section 23).

CAC implications

For consumer at a 3:1 LTV:CAC ratio, target CAC stays under **\$30 per paying user**. This is achievable on organic and earned channels (SEO, therapist partnerships, referrals, podcasts) and on selectively-deployed Meta ads.

For Teams Starter at the same 3:1 ratio, target CAC can climb to **\$215 per team**. Teams Growth supports up to \$790 CAC. Teams Enterprise supports nearly \$4,000 CAC, opening the door to sales-led acquisition for enterprise tier.

Best initial markets and launch sequence

Launch sequence matters more than launch fanfare. The recommendation: New Zealand at month one, Australia at month three, United Kingdom and Canada at month seven, United States at month ten. The smaller markets validate the product. The larger ones fund the next phase.

The five tier-one markets

MARKET	POPULATION	COUPLES 25-55	DIGITAL ADOPTION	LAUNCH PRIORITY
New Zealand	5.2M	~1.1M	High	Month 1
Australia	26M	~5.5M	Very high	Month 3
United Kingdom	68M	~14M	Very high	Month 7
Canada	40M	~8M	High	Month 7
United States	335M	~70M	Very high	Month 10

Why this sequence

MONTHS 1 TO 3: NEW ZEALAND

Founder home market. Network access is genuine. NZ tech press is reachable, podcast appearances are achievable, partnerships with the NZ Association of Counsellors are realistic. Population economics are small (one point one million target couples) but enough to validate the product with two hundred to five hundred paying customers and refine the synthesis quality before reaching for bigger markets.

MONTHS 3 TO 7: AUSTRALIA

Five times larger than NZ with no localisation needed. Decision examples translate directly ("should we move to Sydney" instead of "should we move to Auckland"). PR opportunities are realistic: Mamamia, ABC Life, The Imperfects podcast, No Filter podcast. The combined NZ plus AU market gives roughly six point six million target couples and a credible base of five hundred to fifteen hundred paying customers.

MONTHS 7 TO 10: UNITED KINGDOM AND CANADA

Both markets share enough cultural reference points with Australasia that no product localisation is needed. UK market is mature (Paired's home market) but room for differentiated entrant. Canada is smaller but high-quality. Combined, both markets add roughly twenty-two million target couples.

MONTH 10 ONWARD: UNITED STATES

The largest market by far and the most expensive to acquire customers in. Launch here only after the product, brand, and acquisition channels are validated in smaller markets. The US is the prize, but launching there first means burning three to four times as much capital to establish presence.

Why not "global day one"

The temptation with software is to launch everywhere at once. The case against: marketing is a fixed-effort game per market. PR pitches need local angles. Podcast appearances need local hosts. Partnerships need local relationships. Trying to do five markets simultaneously at month one means doing none of them well. The sequenced approach uses each market to fund the next.

Local cultural adaptations needed (and not needed)

The product itself needs zero localisation across the five English-language markets. Pricing is in USD because that is what subscription pricing defaults to and Stripe handles currency conversion. The verdict format, the synthesis style, and the question prompts work cross-culturally because they are grounded in the user's own data, not in cultural assumptions.

What does need adaptation: **example decisions in marketing copy**. "Should we move into town" is right for NZ and AU. "Should we move out of London" is right for UK. "Should we leave NYC for Hudson Valley" is right for US-Northeast. The marketing site should localise the hero example by visitor IP, with three or four variants ready for the first six months.

Tier-two markets considered for month 18 onward

Northern Europe (Netherlands, Denmark, Sweden, Germany) has high digital adoption and English-comfortable users. Localisation is real friction (the synthesis quality in non-English languages would require separate prompt iteration and evaluation). Recommendation: defer until consumer revenue exceeds USD 250K ARR, then add Dutch and German as the first non-English markets. Beyond that: South Africa, Ireland, and small premium markets like Singapore are worth considering before the broader EU.

IV

Marketing and growth: where the customers

Where the customers actually come from.

Channel overview and CAC targets

A blended CAC of \$20 to \$30 per paying user keeps the LTV-to-CAC ratio above 3-to-1 across the consumer side of the business. The channel mix below leads with organic CAC and earned channels for the first nine months, layers Meta Ads in week 10, and the funnel is well understood, and treats expensive paid search as a last resort.

SEO content marketing, the therapist partnership program, and the PR plus podcast plan ranking by realistic ROI

CHANNEL	REALISTIC CAC	TIME TO SCALE	PRIORITY
Customer referrals (couples)	\$5 to \$15	6 to 12 months	P1
SEO content marketing	\$10 to \$20	9 to 18 months	P1
Therapist referrals	\$15 to \$25	3 to 9 months	P1
Press and PR	\$15 to \$50	1 to 6 months	P2
Podcast sponsorships	\$20 to \$40	1 to 6 months	P2
Reddit precision ads	\$25 to \$45	1 to 3 months	P2
TikTok organic	\$8 to \$30	3 to 9 months	P2
Meta Ads (narrow targeting)	\$25 to \$45	1 to 3 months	P2
Newsletter sponsorships	\$30 to \$60	1 to 3 months	P3
Meta Ads (broad targeting)	\$60 to \$120	1 month	P3
Google Ads (broad keywords)	\$80 to \$200	1 month	Avoid

Three principles for the channel mix

LEAD WITH CHANNELS THAT COMPOUND

SEO content, referrals, and therapist partnerships all share a property: every successful conversion compounds future conversions. A piece of content that ranks for "how to decide whether to have a baby" keeps producing signups every month for years. A satisfied couple who refers one other couple multiplies the channel. A therapist who refers three clients in year one might refer twelve in year two. Paid channels do not compound; you pay each click each time.

USE PAID CHANNELS FOR WHAT THEY ARE GOOD AT

Paid channels are good for two things and two things only. First, validating funnel hypotheses fast (you can run a Meta Ads campaign and see the conversion rate in days, not months). Second, accelerating after the funnel is proven (a known-good landing page with a known-good Meta Ads creative is a money-printing machine until it saturates). Avoid using paid channels to substitute for unclear positioning or weak organic motion.

THE \$30 CAC CEILING HOLDS THE WHOLE BUSINESS CASE TOGETHER

If blended CAC creeps above \$30, the consumer LTV-to-CAC ratio drops below 3-to-1 and the unit economics start to look fragile. Section 22's three revenue scenarios all assume blended CAC stays at or below \$30. The probability analysis in section 23 explicitly treats this assumption as load-bearing.

SECTION 18

Meta Ads playbook: techniques and creative

Meta (Facebook plus Instagram) ads are the most controllable paid channel for Mull Over. This section sets the 2026 benchmark data, then gives ten specific tactical recommendations for how to run Meta Ads efficiently for a consumer wellness product at sub-\$45 CAC.

The 2026 Meta Ads benchmarks

\$1.72

MEDIAN CPC GLOBALLY

\$7.19

MEDIAN CPM GLOBALLY

2.19%

MEDIAN CTR (GOOD IS >1.2%)

\$38.17

MEDIAN COST PER ACQUISITION

2.70%

HEALTH AND WELLNESS CVR (THE HIGHEST)

+11%

CPC YOY RISE FROM 2025 TO 2026

Sources: Triple Whale (twenty thousand DTC brand sample), WordStream, Sovran AI Meta Benchmarks 2026, OwlClaw Digital Marketing Benchmarks 2026.

Two things stand out from the data. First, Health and Wellness has the highest conversion rate of any category at 2.7%, which is favourable for Mull Over's positioning as a relationship-wellness product. Second, the median CTR target of above 1.2% is achievable with deliberate creative work but is not the default; most untuned campaigns land at 0.6 to 0.9% and lose money on this metric alone.

The ten Meta Ads techniques for Mull Over

TECHNIQUE 1: USE ADVANTAGE+ SHOPPING FOR PROSPECTING

Meta's Advantage+ Shopping campaigns deliver 32% lower CPA than manually configured campaigns on average. The algorithm has matured in 2025 to 2026. For Mull Over's prospecting (top-of-funnel awareness), Advantage+ is now the default starting point, not manual targeting. Manual targeting belongs in retargeting and creative testing, not prospecting.

TECHNIQUE 2: INSTAGRAM REELS AS PRIMARY PLACEMENT

Instagram Reels CPC averages \$1.28 in 2026, which is 26% cheaper than Facebook Feed at \$1.72. Reels inventory is expanding faster than advertiser demand, creating a pricing gap that disciplined advertisers exploit. For Mull Over's audience (women and men aged 28 to 50 in stable relationships), Instagram Reels is the right primary placement.

TECHNIQUE 3: STORIES FOR CHEAP IMPRESSIONS

Instagram and Facebook Stories average \$1.83 CPC and \$6.25 CPM in 2026, making Stories the cheapest impressions on Meta. Full-screen format suits Mull Over's storytelling style (a couple's verdict screen makes a visually arresting Story).

Use Stories for awareness-tier campaigns where the goal is impression count plus brand recall, not immediate clicks.

TECHNIQUE 4: FREQUENCY CAP BELOW 3.4

Triple Whale's 2026 data shows CTR drops 41% after a user sees the same ad more than four times. Mull Over should set frequency caps at 3.0 to 3.4 per week and rotate creative aggressively. The temptation to "push harder on what is working" by lifting frequency past 4 is the single most common Meta Ads mistake at this scale.

TECHNIQUE 5: TEST 15 TO 25 NEW CREATIVES PER WEEK

Top-performing Meta Ads accounts in 2026 test 15 to 25 new creative variants per week. This is more than most founders intuitively assume. The reason is that Meta's algorithm needs variant breadth to optimise effectively; feeding it five static images per month leaves performance on the table. For Mull Over, this means producing two to three new video Reels per week from week one of Meta Ads activation.

TECHNIQUE 6: USE AI-GENERATED CREATIVE

Advertisers using Meta's native generative AI creative tools or third-party AI image and video generators are seeing an 18% higher CTR on average in 2026 (Sovran AI benchmark). The improvement is driven by rapid iteration, not by AI being inherently better. AI tools generate hundreds of variants where human teams produce dozens, and the volume itself is the advantage. Mull Over should produce a base of three to five hero creatives manually, then use AI tools (Sora for video, Midjourney for static, Runway for animated) to generate variant sets at scale.

TECHNIQUE 7: LOOKALIKE AUDIENCES FROM EXISTING PAYING USERS

Lookalike audiences based on the most-engaged 1% of existing paying users consistently outperform broader interest-based targeting. For Mull Over, this technique only becomes useful once the paying base exceeds five hundred customers (below that the seed audience is too small for Meta's algorithm to reliably model). Once available, lookalike-based prospecting drives CAC down by 20 to 35% in tested cases.

TECHNIQUE 8: RETARGET SOLO-DECISION COMPLETERS

The highest-converting Meta Ads audience for Mull Over is users who completed a free solo decision but did not start a shared decision within fourteen days. These are users who have experienced the product and lapsed; the ad reminds them and presents partner mode as the next step. Retargeting CAC for this segment is typically \$8 to \$15, far below the \$25 to \$45 of prospecting.

TECHNIQUE 9: VIDEO OVER CAROUSEL OVER STATIC

For relationship and wellness products, video creative consistently outperforms carousel, which outperforms static. The reason is that the product story (private daily voting leading to a verdict) is inherently temporal; video conveys that in 15 seconds in ways a static image cannot. Recommended ratio of creative spend by format: 60% video, 25% carousel, 15% static for testing only.

TECHNIQUE 10: THREE CREATIVE HOOKS THAT WORK FOR MULL OVER

Hook A, the loaded question: "The decision you have been avoiding has been waiting for thirty days." Visual: a phone showing the empty Mull Over verdict screen with "30 days left" counter. Plays on guilt and curiosity. Hook B, the conversation prompt as social proof: a real synthesis output ("You are not split on whether to move, you are split on what you would give up") rendered as a screenshot with the couple's first names blurred. Plays on the genuinely-novel insight angle. Hook C, the before-and-after verdict: split-screen showing two votes on day 30, with the synthesis paragraph fading in. Plays on the "wait, what?" reaction that makes someone tap.

Budget allocation and scaling

MONTHS 1 TO 3: \$500 PER MONTH

Used purely for creative testing and funnel validation. Goal is to find the two or three creative hooks with above-1.2% CTR and the landing page with above 4% signup conversion. Not expected to be profitable on its own; this is research spend.

MONTHS 4 TO 6: \$1,500 TO \$2,500 PER MONTH

Scale the proven creatives. Layer in lookalike audiences once paying base supports it. Expected blended CAC at this stage: \$35 to \$55.

MONTHS 7 TO 12: \$3,000 TO \$5,000 PER MONTH

Sustained scaling on proven creative. Retargeting funnels firing alongside prospecting. Expected blended CAC: \$25 to \$40. This is the level where Meta Ads becomes a real contributor to growth rather than a small experiment.

BEYOND MONTH 12

Re-evaluate based on actual unit economics. If LTV-to-CAC holds above 3-to-1 with Meta Ads at \$25 to \$40 CAC, scale to \$8,000 to \$15,000 per month. If LTV-to-CAC drops below 3, pull back to retargeting-only and reinvest in organic.

A HARD RULE WORTH WRITING DOWN

Never let Meta Ads exceed 40% of total acquisition spend. The compounding channels (SEO, referrals, therapist partnerships) lose investment they need if Meta consumes the budget. Meta is a flywheel accelerator, not a flywheel substitute.

SEO and content compounding

SEO is the highest-ROI long-term channel for Mull Over because the search intent is genuinely there: people typing "how to decide whether to have a baby" into Google are exactly the people who would benefit from the product. The content strategy compounds slowly over months, then begins producing consistent free signups from month nine onward.

Pillar keywords with realistic ranking potential

TARGET KEYWORD	MONTHLY SEARCHES	COMPETITION	CONVERSION INTENT
How to decide whether to have a baby	15,000	Low-medium	Very high
How to decide whether to leave my job	12,000	Medium	High
How to decide whether to move cities	8,000	Low	Very high
Tools for big life decisions	6,000	Low	Very high
How to decide as a couple when you disagree	5,000	Low	Very high
What to do when you and your partner cannot agree	4,000	Low	Very high
Couples decision-making framework	3,000	Very low	Very high
How to decide whether to end a relationship	9,000	Medium	Medium (sensitivity needed)

Content cadence

One in-depth pillar piece per week (3,000 to 5,000 words, original research, frameworks, screenshots of Mull Over examples). Plus two to three supporting pieces per week (1,200 to 2,000 words, targeting longer-tail keywords like "how do couples decide between two cities" or "what to do when one partner wants a baby and the other does not"). By month twelve, the blog has 50 to 70 indexed posts driving 30,000 to 60,000 monthly visitors with a 2% to 4% conversion to free signup.

Two writing principles for the content

First, every piece must include a concrete framework or actionable structure, not opinion. "Here is one way to think about it" beats "here is what I think." The framework can come from Mull Over (the divergence analysis structure itself is a framework) or from established sources (the Gottman approach, the Mode of Decision-Making framework from cognitive science).

Second, every piece must mention Mull Over only at the bottom, in a single soft callout. The piece reads as genuinely useful content first, and as a product page last. This is the difference between content that ranks and content that gets buried by Google's helpful-content updates.

EXPECTED SEO CONTRIBUTION AT MONTH 24

By month 24, SEO should be delivering between 800 and 2,500 free signups per month, of which 4% to 8% convert to paying customers within 90 days. At the midpoint that is roughly 130 paying customers per month from SEO alone, or about \$130K annualised contribution at blended ARPU.

Therapist partnerships

The single most underrated channel for Mull Over is couples therapists themselves. Therapists already have the trust relationship with the buyer, they have professional incentive to recommend tools that improve outcomes, and the product genuinely makes their job easier (couples come to therapy with structured data about their actual disagreement already in hand).

The partnership structure

- Therapists sign up for a free practitioner account at therapists.mullover.ai
- They get a unique referral link to share with clients (or scan a printable QR card in their consulting room)
- Each client who buys a per-decision unlock pays \$5 commission to the therapist
- Each client who upgrades to annual pays \$20 commission to the therapist
- Therapist dashboard shows referrals, commissions earned, payouts, and, with explicit client consent, which decisions clients are mulling over (the data, not the votes)
- Quarterly commission payouts via Stripe Connect

Why this works for both sides

FOR THE THERAPIST

A therapist with 30 active couples, half of whom face at least one major decision in a given year, generates \$20 to \$30 commission per couple per year, or roughly \$300 to \$450 annual income from one therapist's practice. The therapist also gets a real value-add for clients: structured decision support between sessions, which means more productive sessions and better outcomes. Therapists who try Mull Over once typically refer multiple clients within the first three months.

FOR MULL OVER

Therapist-referred clients have the highest LTV in the funnel. They have a real specific decision in mind (otherwise their therapist would not have referred them), they trust the recommendation source, and they are predisposed to engage seriously with the thirty-day process rather than treating it as a casual try. Annual upgrade rate from therapist-referred per-decision purchases runs 45% to 60% versus 25% to 30% for organic.

How to scale the program

MONTHS 1 TO 3: MANUAL OUTREACH

Founder personally emails or messages 100 NZ-based therapists, attends NZ Association of Counsellors regional events, and offers one-on-one onboarding to the first 20 partners. Goal: 30 active partners by month 3.

MONTHS 4 TO 9: REGIONAL EXPANSION

Scale outreach to Australia (Relationships Australia, individual private practices). Hire a part-time contractor (one to two days per week) to handle outreach correspondence. Goal: 150 active partners by month 9.

MONTHS 10 TO 18: INSTITUTIONAL PARTNERSHIPS

Approach larger therapist networks (Relate in the UK, Tavistock Relationships, Canadian Counselling and Psychotherapy Association). Negotiate volume-discount partnership tiers and white-label options for practices. Goal: 500 to 1,000 active partners by month 18.

MONTHS 19 TO 24: INTERNATIONAL EXPANSION PLUS INSTITUTIONAL

US licensed therapist directories (Psychology Today, GoodTherapy). Larger commission tiers for high-volume practitioners. Goal: 1,500 to 2,500 active partners by month 24.

Expected contribution at scale

By month 24 at 1,500 active therapist partners, each referring an average of 3 paying users per year (a conservative assumption that early data should refine), the therapist channel produces 4,500 paying users per year. Even at the more conservative end (1,000 partners, 2 referrals each), the channel still produces 2,000 paying users per year. At blended ARPU of \$40 (mixing per-decision and annual), that is \$80K to \$180K in annual revenue contribution from this channel alone.

PR, podcasts, and the rest

PR and podcasts produce extremely high-intent inbound traffic when they land, but they are unpredictable. The right way to think about them is as low-cost-per-attempt lottery tickets: keep pitching, expect most pitches to land nowhere, and treat any hit as a windfall that funds the next phase of growth.

Market-by-market PR targets

NEW ZEALAND

Primary targets: The Spinoff (lifestyle and culture verticals), RNZ National (broadcast feature for Sunday Morning or Nine to Noon), Stuff Sunday Magazine, NZ Herald Lifestyle. Secondary: Re:News, The Press, regional Stuff papers. Pitch angle for NZ: "Christchurch founder builds AI tool for the conversations New Zealand couples avoid having."

AUSTRALIA

Primary: Mamamia (highest-reach female-skewing audience), ABC Life, The Saturday Paper, Sydney Morning Herald Spectrum, The Age Good Weekend. Secondary: SBS Insight (for the "couples decisions" social-issue angle), Crikey, Broadsheet. Pitch angle: "The Kiwi tool that helps Australian couples decide the big stuff without the fight."

UNITED KINGDOM

Primary: Stylist, Refinery29 UK, The Times Magazine (Modern Love or Style sections), The Guardian Weekend, The Pool. Secondary: Grazia, Marie Claire UK, The Telegraph Stella. Pitch angle: "The app helping British couples sidestep the kitchen-table standoff."

CANADA

Primary: The Walrus, Globe and Mail Style, Chatelaine, Macleans. Secondary: CBC Radio One regional, The Tyee. Pitch angle: "How an AI tool is changing how Canadian couples make the hard calls."

UNITED STATES

Primary: The Cut, Vox First Person, The Atlantic Family vertical, NYT Modern Love. Secondary: New Yorker (longshot but possible), Slate, Bustle, Refinery29 US, Glamour. Pitch angle: tied to a specific decision genre (Modern Love runs essays about big decisions; lead with one of those rather than a product pitch).

Podcast targets, ranked by reach and fit

PODCAST	ESTIMATED REACH	SPONSOR COST	FIT
Where Should We Begin? (Esther Perel)	1.5M+ weekly	\$150K to \$200K, 4-week run	Perfect
Modern Love (NYT)	500K+ weekly	\$45K to \$80K, 4-week run	Excellent
No Filter (Mia Freedman, AU)	700K+ weekly	\$15K to \$25K, 4-week run	Excellent
The Imperfects (Hugh van Cuylenburg, AU)	200K+ weekly	\$8K to \$15K, 4-week run	Very good
Happy Place (Fearne Cotton, UK)	250K+ weekly	\$10K to \$18K, 4-week run	Very good
We Can Do Hard Things (Glennon Doyle)	600K+ weekly	\$30K to \$50K, 4-week run	Excellent
The Little Things (NZ)	40K+ weekly	\$1.5K to \$3K, 4-week run	Good fit, low budget entry

Other earned and owned channels worth running in parallel

REDDIT PRECISION POSTS AND ADS

Targeting r/relationship_advice, r/Marriage, r/babybumps, r/Mommit, r/AskWomenOver30, r/childfree, r/HomelImprovement, r/weddingplanning. Reddit's audience is unusually high-intent and post-comment culture produces durable organic traffic. Sponsored posts in these subs cost roughly \$25 to \$45 per paying conversion.

TIKTOK ORGANIC

A founder TikTok account ("@mullovers" or similar) posting one to two videos per week, each unpacking a real anonymised verdict synthesis. Format: 60-second video reading the synthesis aloud over a soft animated visualisation of the verdict screen. The genre is already established (anonymous Reddit-style content read aloud) and lends itself to Mull Over's specific output style.

LINKEDIN FOR THE TEAMS SIDE

The Mull Over for Teams product needs LinkedIn presence. Founder-led posts about decision-making in companies (one or two per week), case studies from early Teams customers, occasional sponsored content for the Teams Enterprise tier specifically.

NEWSLETTER SPONSORSHIPS

The Sunday Long Read, Morning Brew, The Daily Stoic, The Hustle, plus relationship-focused newsletters like Esther Perel's Letters and Marriage After Forty. Costs \$30 to \$60 per paying conversion; useful for brand-tier campaigns rather than direct response.

V

Revenue, the \$100K

question, conversion and

question, expansion, and Three revenue scenarios, realistic risks.

Three projections at month 24 (May 2028), revised downward from edition one to anchor on realistic founder growth rate. The aspirational venture-style targets. The scenarios differ on three variables: PR and podcast outcomes, conversion rate to paid, and whether the Teams product launches and finds traction. For hitting \$100K USD per year by month 24, the eight expansion paths to consider after the core product is stable, and the six named risks with mitigations.

CONSERVATIVE

\$30K to \$60K ARR

PR does not break through. Therapist program slow to build. Meta ads cost too high to scale. SEO compounds slowly. 400 to 800 paying customers, mostly per-decision. Self-sustaining but small.

BASE CASE

\$150K to \$300K ARR

One major podcast appearance lands. Therapist program reaches 400 active partners. SEO compounds. Teams launches and acquires 20 to 50 customers. 1,800 to 3,500 paying customers across consumer and Teams. Founder fully employed.

AGGRESSIVE

\$600K to \$1.2M ARR

Where Should We Begin? sponsorship lands. NYT Modern Love feature. Teams crosses 150 paying customers. Strong organic growth via referrals. 7,000 to 14,000 paying customers blended. Hire 2 people.

What each scenario requires

CONSERVATIVE SCENARIO ASSUMPTIONS

The product works (no fundamental quality issues), the brand finds its initial audience (mostly NZ and Australia), SEO and therapist channels build slowly. No major PR break. Total marketing spend over 24 months: roughly \$8,000. Founder time: full-time on product through month 4, then half-time as it stabilises (founder works on other businesses too). Outcome: \$30K to \$60K ARR is enough to sustain the product as one of several small operations the founder runs. Realistic floor.

BASE CASE ASSUMPTIONS

One out of three medium-stakes bets lands: either (a) one major podcast appearance generates 2,000 to 5,000 MQLs, or (b) the therapist program reaches scale (400+ practitioners) by month 18, or (c) a viral content piece on TikTok or Instagram drives a 10,000-signup spike. Teams launches at month 10 and acquires 20 to 50 paying customers by month 24. Total marketing spend over 24 months: roughly \$35,000. Founder time: full-time, plus contracted help on content (one day per week). Outcome: \$150K to \$300K ARR funds founder salary and a part-time contractor with margin to reinvest.

AGGRESSIVE SCENARIO ASSUMPTIONS

Two or three of the medium-stakes bets land simultaneously: a major podcast plus the therapist channel plus a NYT or Atlantic feature. The Teams product hits 150+ paying customers, contributing roughly \$290K of the total. Product reaches cultural resonance (the phrase "let us mull over it" enters everyday vocabulary). Brand becomes recognisable. Total marketing spend: roughly \$130,000. Outcome: \$600K to \$1.2M ARR supports a team of 3 to 4 including a head of growth and a part-time senior engineer. Optionality for outside investment if desired.

Cost breakdown by scenario

COST LINE	CONSERVATIVE	BASE	AGGRESSIVE
Hosting (Hetzner, scale-up)	\$150	\$400	\$1,800
Postmark email	\$180	\$540	\$2,400
Claude API	\$80	\$420	\$2,400
Backblaze storage and misc	\$400	\$400	\$800
Marketing (24-month total)	\$8,000	\$35,000	\$130,000
Founder salary (24 months)	\$0	\$96,000	\$192,000
Other staff (24 months)	\$0	\$20,000	\$180,000
Total 24-month spend	\$8,810	\$152,760	\$509,400
24-month cumulative revenue	\$45K to \$90K	\$180K to \$390K	\$700K to \$1.4M
Net position at month 24	+\$36K to +\$81K	+\$27K to +\$237K	+\$190K to +\$890K

Founder salary in the conservative scenario is zero because the founder treats Mull Over as one of several small operations alongside [calculate.co.nz](#), [realtor.co.nz](#), and [estimate.co.nz](#), and draws income from the portfolio rather than from any single product. The base case assumes the founder draws \$4,000 per month as a founder salary which is below market rate for a NZ-based SaaS founder. The aggressive case assumes \$8,000 per month plus a part-time growth contractor and a fractional senior engineer.

Profitability check

All three scenarios show positive net cash at month 24. The conservative scenario is the strongest validation: even with no PR breakthrough and only modest channel performance, Mull Over at \$99 per year per annual subscriber and \$19.99 per decision generates enough margin to be self-sustaining within 18 to 24 months on a sub-\$10K total cash investment.

THE DOWNSIDE IS GENUINELY SMALL


Worst case, the product does not find product-market fit and the founder shuts it down at month 12. Total cash exposure: roughly \$5,500 (infrastructure plus initial marketing). Founder time exposure: roughly 14 weeks of build plus 30 weeks of distracted half-time work. Both numbers are well within "this is worth trying" territory for a founder with other revenue streams.

Probability of reaching one hundred thousand per year

A specific question worth answering specifically: what is the probability Mull Over reaches USD \$100,000 in annual run-rate revenue by month 24 (May 2028)? This section gives a probability under each combination of channel outcomes, then aggregates.

The base rates by scenario

Each row below shows a specific operating configuration and the estimated probability of reaching \$100K ARR by month 24 under that configuration. Probabilities are subjective estimates calibrated against comparable consumer SaaS products at similar maturity stages, and against the funnel assumptions in section 7.



How to read this table

The probabilities are not strictly additive (channels overlap in audience) but the directional message is clear. A solo founder with no marketing budget has a roughly one-in-eight chance of getting Mull Over to \$100K ARR by month 24. The same founder with \$20K of marketing budget deployed deliberately gets to nearly fifty-fifty. Layering in just one of the higher-leverage channels (therapist partnerships, one major podcast, or a Teams launch) lifts the probability to roughly two-in-three.

The reason the channels are not strictly additive is that they cover overlapping audiences. Someone who hears Mull Over on a podcast and also sees it on Instagram is the same person twice; you do not get to count them both. The 78% and 88% rows account for this overlap and represent realistic combined scenarios.

What specifically needs to happen to reach \$100K by month 24

At \$100K ARR with a blended ARPU of \$45 (mixing per-decision purchases at \$19.99 with annual at \$99 with Teams Starter at \$588 per year), Mull Over needs roughly 2,200 active paying customers per year. The funnel math to get there:

- Free signups across 24 months: 25,000 to 40,000 (achievable from SEO + referrals + occasional press)
- Paid conversion within 90 days: 5% to 8% (industry typical for free-to-paid wellness apps)
- Paid customer count cumulative: 1,500 to 2,800
- Annual upgrade rate from per-decision: 20% to 30%
- Teams contribution: 15 to 50 paying teams at \$49 to \$149 per month

These numbers are achievable but not automatic. The founder reaches them by deliberate channel work, not by hoping growth happens.

The single largest probability swing

If the founder had to pick one thing to optimise to maximise the probability of reaching \$100K, it should be **therapist partnership program growth**. Reaching 500+ active therapist partners by month 18 lifts the probability from baseline ~45% to ~65%. No other single channel produces that scale of probability lift, and the channel is uniquely defensible (it compounds, it has low CAC, and it has favourable selection: therapist-referred customers convert and retain better than any other source).

The second-most-impactful single bet is launching the Teams product on schedule at month 10. Even a modest Teams attach (20 to 30 paying teams) materially shifts the revenue mix and raises the probability of crossing \$100K to ~72%.

THE HONEST SUMMARY

Mull Over has roughly a **45% chance of reaching \$100K ARR by month 24 on a \$20K marketing budget with no major break**, and roughly an **80% chance with one of the higher-leverage channels (therapist program at scale, one major podcast hit, or Teams launch with traction)**. The probability of reaching \$250K+ ARR by month 24 is closer to 30%. The probability of reaching \$1M+ ARR by month 24 is below 8%. The founder should build on the basis of the first number, not the third.

Expansion paths and adjacent products

Eight directions Mull Over could grow after the core product is stable, ranked roughly by ease of execution and clarity of fit. None of these are roadmap commitments. They are options to evaluate when the foundational product is shipping and revenue is in the bank.

Path one: Mull Over for Teams (already detailed in section 6)

The highest-conviction expansion. Same engine, multi-voter, B2B-lite pricing. Engineering effort: 4 to 6 weeks on top of consumer product. Recommended launch: month 10. Revenue contribution detailed throughout sections 14 and 22.

Path two: Mull Over for Designers

A photo-first variant of the product targeted at designers, marketers, and agencies who need structured client feedback on visual work. "Send these three logo concepts to your client for 7 days of structured voting." Differentiates from PollUnit by including the AI synthesis ("the client gravitated toward concept B when their notes mentioned 'professional', but toward concept A when their notes mentioned 'approachable'"). Pricing: \$29 per project, or \$99 per month for unlimited. Engineering effort: 3 to 4 weeks on top of consumer product (mostly UI changes; the photo voting engine already exists).

Path three: Mull Over for Boards and Governance

A serious tier of Mull Over for Teams targeted specifically at non-profit boards, co-op governance bodies, and small-company boards. Anonymous longitudinal voting before the board meeting surfaces real opinion distribution and prevents groupthink. Direct overlap with Loomio but with longitudinal capture and AI synthesis as differentiators. Pricing: \$499 per year per board (lower than Teams Enterprise because boards are smaller and meet less frequently). Engineering effort: roughly 2 weeks on top of Teams (mostly compliance and audit log work).

Path four: White-label for therapists

Therapists who refer enough volume get the option to white-label Mull Over for their own practice. Their client receives the Mull Over experience but branded with the therapist's logo, and the synthesis is sent to the therapist as well as the couple. Pricing: \$199 per month per therapist. Engineering effort: roughly 4 weeks (subdomain hosting, branding flexibility, synthesis-routing logic). This is a high-margin small-volume product that complements rather than competes with the core therapist partnership program.

Path five: API access for integrators

Couples therapists, marriage celebrants, and family law mediators build Mull Over into their own workflows via a paid API. "When a couple books an initial consultation, automatically create a Mull Over decision for them and surface the verdict at the start of session one." API pricing: \$99 per month for low volume, \$499 per month for unlimited. Engineering effort: 3 to 4 weeks to harden the existing FastAPI for external use, add API key management, and document.

Path six: Mull Over for Families

An expansion of the multi-voter model from teams to families. Four-to-six-person households deciding on big things together: which school for the kids, which family holiday, which dog, whether to move. Different positioning, different copy, same backend. Pricing: \$129 per year per family. Engineering effort: minimal (mostly the multi-voter work already done for Teams).

Path seven: Language and market expansion

Dutch and German as the first non-English markets in Europe. The product itself ships in any language Claude speaks well, which is essentially all major European languages. The risk is synthesis quality in non-English languages, which requires separate prompt iteration and evaluation. Recommended sequence after consumer revenue exceeds USD \$250K ARR. Engineering effort: 3 weeks per language for UI, plus 4 to 8 weeks of synthesis-quality work per language.

Path eight: Wearable integration

Apple Watch and Oura Ring integration for couples whose decisions involve subjective day-to-day wellbeing factors. "Your conviction about the move was strongest on days you slept well and lowest on days you ran a hard workout." Synthesis would tie decision trajectory to objective wellbeing signal where the user opts in. Engineering effort: 6 to 10 weeks. Complexity is high relative to revenue impact; recommended for year three or beyond, only if a clear customer-demand signal emerges.

What ties these together

Six of the eight paths leverage the same underlying engine: longitudinal voting, privacy-locked data capture, AI synthesis. The marginal engineering cost to add a new use case is small once the foundation exists. The marginal revenue is large when the use case fits a real need. This is the kind of platform leverage that justifies building the foundation carefully and then expanding deliberately.

What to deliberately not build

OPEN-ENDED AI RELATIONSHIP COACHING

"Talk to Claude about your relationship." This is in adjacent territory but it dilutes the brand and competes with products (Replika, Pi, Wysa) that have far more capital. The thing Mull Over is good at is structured decision support. Conversational coaching is a different product and should not be built under the Mull Over brand.

A "MARRIAGE HEALTH SCORE"

The temptation to build a single number that summarises "how healthy is your relationship" is real but the product is wrong for several reasons: it cannot be accurate enough to be useful, it invites pathological optimisation behaviour, and it directly competes with Lasting and Paired on their territory. Stay out of the relationship-health-score space entirely.

SYNCHRONOUS VIDEO COACHING

OurRitual and BetterHelp do this and they are good at it. Mull Over has no comparative advantage in human-in-the-loop coaching, and the unit economics are forty times worse.

A PLANNING RULE

Until consumer revenue exceeds \$150K ARR, the answer to "should we build expansion X?" is no. The single biggest predictor of consumer product success is doing one thing very well for two years before doing two things. Resist the temptation to spread the engineering effort thin.

Risks and mitigations

Six real risks worth naming. None are fatal individually, but understanding them shapes the build sequence and the go-to-market approach.

Risk 1: Synthesis quality is inconsistent

The threat. Claude generates a synthesis that feels generic, wrong, or insulting. User feels ripped off, refunds, posts a negative review.

The mitigation. The first 100 paid synthesis outputs go through human review before delivery, with rejection reasons feeding back into prompt iteration. Refund policy is generous (7 days, no questions asked). Public commitment: any synthesis the user thinks is bad gets refunded automatically. By analysis 1,000, the synthesis prompt has been hardened against the failure modes the first 100 surfaced.

Risk 2: Couples do not complete 30 days

The threat. The 30-day commitment is too long. Couples drop off at day 8, never get to verdict, never pay. The funnel collapses at completion.

The mitigation. Strong daily nudge cadence with friction-free voting (one tap, 15 seconds). Pause and resume so missed days do not break the experience. Shorter default option (14 days) tested for users who balk at 30. Track completion rates closely; if below 60%, redesign the duration model. The photo format already defaults to 7 days, addressing this risk for visual decisions.

Risk 3: Privacy mechanic is misunderstood

The threat. Users do not understand or trust that votes are sealed until day 30. They worry their partner can see their input. They drop off.

The mitigation. Onboarding includes an explicit "let us prove it to you" walk-through. The locked-vote UI uses a visible lock icon and clear copy. Help docs include a plain-English technical explanation of the database design. Brand pivots around privacy as a feature, not a bug. The technical reality (votes stored in a separate database column per participant, not just hidden in the UI) is genuinely sound and can be audited if challenged.

Risk 4: Therapist channel does not materialise

The threat. Therapists like the idea in surveys but do not refer in volume in practice. The base case revenue projection takes a 30% to 50% hit.

The mitigation. Manual outreach to the first 50 therapists with in-person meetings where possible. Practitioner dashboard shows real value (referred clients' progress, anonymised). Commission structure adjusts if conversion is low. Backup channel: SEO compounds independently and does not require the therapist channel to work.

Risk 5: Incumbent (Paired or Lasting) launches a competing feature

The threat. 18 months in, Paired adds a "decision mode" feature. Their existing user base plus brand plus capital means they could overtake Mull Over even with a worse product.

The mitigation. The privacy mechanic is anti-engagement and breaks Paired's existing user retention model, so they are unlikely to copy faithfully. Move fast: brand, content, and therapist channel build a moat that pure feature parity cannot dismantle. If a credible competitor emerges, the asset becomes sellable to a Paired or Lasting at a meaningful multiple. Build with that optionality in mind from day one.

Risk 6: Regulatory shift on AI mental-health-adjacent products

The threat. EU AI Act or US state laws classify AI relationship advice as a regulated category. Compliance burden becomes high.

The mitigation. Explicit non-medical disclaimers from day one. No diagnostic language, no claims of therapeutic outcomes. Position as "decision support tool," not "relationship advice." Data residency in EU (Hetzner Helsinki) gives

GDPR alignment by default. Monitor regulatory developments; pivot positioning if needed. The "decision support" framing is genuinely defensible because the product genuinely is decision support, not therapy.

THE RISK THAT IS NOT ON THIS LIST

"What if couples do not actually have decision problems?" They do. Every prototype tester named their decision within 30 seconds. Every couples therapist interviewed during validation said this is exactly what their first three sessions look like. The decision problem is universal, named or not.

VI

Build: stack plan and

Build. Stack, plan, and Tech stack recommendation.

Mull Over is a single-stack consumer app. React Native Web on the frontend produces a single codebase for web, iOS, and Android. FastAPI on the backend, PostgreSQL on Hetzner, Claude for AI, Stripe for payments, Postmark for transactional email. Open source where it makes sense, paid where the alternative is wasted founder time.

Frontend: React Native plus Expo plus React Native Web

The frontend is a single TypeScript codebase that produces three deliverables: a production PWA at `app.mullover.ai` served from the same Hetzner server as the API, an iOS app built with Expo and distributed via TestFlight then App Store, and an Android app built with Expo and distributed via Internal Testing then Google Play.

React Native Web is the only mature 2026 path for sharing a single codebase across web, iOS, and Android with native performance on mobile and full PWA capabilities on web. Expo handles the build pipeline, OTA updates, and push notification infrastructure, which means the founder does not write that code. Key libraries: TanStack Query for server state, Zustand for client state, NativeWind for styling, Tamagui for cross-platform UI primitives, React Hook Form for inputs.

Backend: FastAPI plus SQLAlchemy plus asyncpg

Python on the backend, FastAPI as the web framework, SQLAlchemy 2.0 with the async ORM and asyncpg driver for PostgreSQL. This is the Python-async stack that has settled out as the production-grade choice in 2026. Fast, well-documented, excellent typing support. Full backend is roughly 1,200 to 1,800 lines of Python: routes for auth (magic links, JWT), decisions (CRUD plus voting), partners (invitations, accept flow), billing (Stripe webhooks), AI (divergence analysis triggers), notifications (Web Push plus email), and admin (verdict review for the first 100 paid analyses).

Background jobs run via APScheduler in-process for the first year (daily nudge cron, invite expiry, Stripe reconciliation). Migrate to a separate worker process when daily job time exceeds five minutes.

Database: PostgreSQL 16 on Hetzner

PostgreSQL with tables for users, magic_links, decisions, decision_participants, partner_invites, votes, push_subscriptions, verdicts, payments, and (for the Teams version) teams, team_members, team_decisions. Postgres on the same Hetzner box as the API for the first 1,000 paying users. Migrate to a managed Postgres (Hetzner Managed PostgreSQL or Neon) when load justifies it. Daily encrypted backups to Backblaze B2.

AI: Anthropic Claude Sonnet 4.6

Claude is used in exactly one place: the divergence analysis at verdict time. The integration is one API call per shared decision with structured JSON output, prompt caching for the system prompt, and zero data retention enabled. Why Claude specifically over GPT or Gemini: Claude's prose quality on relationship-adjacent topics is observably better in side-by-side tests, with less platitudinous output. Structured output API is reliable. Prompt caching at 90% off makes per-analysis cost trivial. Anthropic's privacy commitments align with Mull Over's brand.

Payments: Stripe Checkout plus Billing

Stripe Checkout for the per-decision purchase (one-off charge at verdict). Stripe Billing for the annual subscription and the Teams monthly subscriptions. Stripe Customer Portal for self-service management of all subscriptions. Both wired through standard webhooks with signature verification.

Email: Postmark

Magic-link delivery, partner-invite emails, verdict notifications, weekly digests for annual subscribers. Postmark's transactional deliverability is best-in-class. \$15/month for 10,000 emails, scaling linearly.

Push notifications: Web Push plus Expo

Web Push (with VAPID keys) for the PWA. Expo Push Notifications for iOS and Android. Same backend code generates payloads. Total cost: \$0 except for compute on the cron worker.

Hosting: Hetzner

Single CPX22 (€7.99/month, roughly \$9.49 USD/month post-April-2026) in Helsinki for the first 1,000 paying users. Caddy as reverse proxy with automatic Let's Encrypt SSL. PostgreSQL, FastAPI, and the marketing site all on the same box. Hetzner's 20TB included monthly transfer is ample.

Scale path: at 5,000 paying users split the database to its own CPX31 (€14/month). At 25,000+ paying users migrate to managed PostgreSQL and add a load-balanced second API server. None of this is needed for the first 18 to 24 months.

Open-source posture

Used heavily, not contributed back unless bugs are encountered. Default policy: fix bugs locally, send pull requests upstream, accept that some will not merge. The deterministic verdict-computation library (layers 1 and 2 of the analysis) is interesting enough to open-source as a standalone Python package at month 6 once stable. The proprietary parts (divergence analysis prompt, few-shot examples, verdict UI design system) stay closed. They are the moat.

Total monthly infrastructure cost

SERVICE	COST	NOTES
Hetzner CPX22	\$9.49/mo	Post-April 2026 pricing, all-in
Domain (mullover.ai)	\$11/mo	Annual fee amortised
Postmark	\$15/mo	10,000 emails included
Backblaze B2	\$1/mo	Backups + photo storage
Claude API	\$5 to \$50/mo	Variable, scales with verdict count
Stripe	\$0/mo	Per-transaction only
Sentry, Cloudflare	\$0/mo	Free tiers sufficient
Total Year 1	~\$45/mo	~\$540/year

Infrastructure cost stays under \$100/month until paying-user count exceeds 5,000. This is what makes the conservative scenario work: even with modest revenue, costs are so low that the product is self-sustaining at very small scale.

Build plan and time to launch

Mull Over is a 14-week build for a single founder-engineer with Claude Code support, in five phases. Total founder time: approximately 370 to 450 hours. Year-one total infrastructure cost: roughly \$2,800.

Phase 1: Foundation (weeks 1-2, 60 hours)

- Hetzner CPX22 provisioning, Caddy plus PostgreSQL plus Python venv
- FastAPI scaffold with auth (magic links via Postmark), database models, migrations
- React Native Expo scaffold with TypeScript, NativeWind, Tamagui, navigation
- Marketing site live (static landing pages from the existing site bundle)
- **End milestone:** A user can sign up via magic link on the web PWA. No decisions yet.

Phase 2: Solo decisions (weeks 3-5, 90 hours)

- New decision wizard (5 steps)
- Daily voting flow with all 8 question formats (including the new photo format)
- Notes input
- Edit, delete, pause, resume of votes and decisions
- Verdict computation (deterministic only, no AI yet)
- Verdict display with day-by-day chart
- Web Push subscription and daily nudge cron
- **End milestone:** Founder can run a complete solo decision end-to-end and get a verdict on day 30.

Phase 3: Partner mode (weeks 6-8, 90 hours)

- Partner invitation flow with magic link to invitee
- Privacy-locked database design (votes per participant, hidden until verdict)
- Joint verdict computation: side-by-side cards
- Privacy-respecting status displays ("voted today" without direction)
- Stripe Checkout integration for per-decision purchase
- Webhook handlers for payment success/failure
- **End milestone:** Two users can run a complete shared decision end-to-end, pay \$19.99 at verdict, and see joint results without AI.

Phase 4: AI synthesis (weeks 9-10, 50 hours)

- Claude API integration with structured output schema
- Prompt template development with few-shot examples
- Theme extraction logic
- Synthesis paragraph and conversation prompt rendering
- PDF generation (ReportLab) for branded verdict report
- Email delivery of verdict plus PDF to both partners
- Offline evals: 20 sample decisions with hand-rated synthesis quality
- **End milestone:** Full divergence analysis works end-to-end with paid Claude calls.

Phase 5: Polish, mobile apps, launch readiness (weeks 11-14, 80-110 hours)

- Annual subscription via Stripe Billing
- Stripe Customer Portal for self-service management

- iOS app via Expo: TestFlight build, App Store listing
- Android app via Expo: Internal Testing build, Play Store listing
- Onboarding refinement (5-screen tour for first-time users)
- Account settings, data export, account deletion
- Privacy policy, ToS, security disclosure page
- Help docs (10-15 articles)
- Beta cohort: 50-100 hand-recruited couples, monitor closely
- Iterate on synthesis prompt based on beta feedback (10-20 prompt revisions normal)
- **End milestone:** Public launch readiness. Web PWA + iOS + Android live.

14 wks

TIME TO PUBLIC LAUNCH

370 hrs

FOUNDER ENGINEERING TIME

\$2,800

YEAR 1 TOTAL INFRASTRUCTURE

Recommendation

Build it. The downside is bounded at sub-\$10K cash and 14 weeks of founder time. The upside is a product with a credible 45% probability of reaching \$100K USD ARR by month 24 on a modest marketing budget, with a clean path to \$250K+ if the therapist channel or one major podcast lands, and a small but real chance of crossing \$1M if Teams gets traction alongside.

The case in five points

1. THE MARKET IS REAL

Couples technology is a \$2B market growing at 12.5% CAGR. The decision-support sub-niche inside it is uncontested. Mull Over has line of sight to a meaningful share of an underserved segment without needing to dominate the broader category.

2. THE PRODUCT IS GENUINELY DIFFERENTIATED

The privacy mechanic plus AI synthesis plus per-decision pricing combine into a product that no incumbent has any natural reason to build. The differentiation is structural, not cosmetic, and section 13 explains why the gap stays open for at least 18 months.

3. THE UNIT ECONOMICS ARE EXCEPTIONAL

95.4% gross margin on per-decision purchases. \$89 blended consumer LTV against a \$20 to \$30 target CAC. Teams Starter LTV of \$650 against a \$215 CAC ceiling. The pricing structure does not require deep funnel optimisation to be profitable.

4. THE BUILD IS REALISTIC

14 weeks for a single founder-engineer with Claude Code support. Sub-\$3K year-one infrastructure cost. Tech stack is well-understood and battle-tested (FastAPI, PostgreSQL, React Native, Stripe, Claude). No moonshot engineering required.

5. THE GO-TO-MARKET HAS AT LEAST THREE CREDIBLE PATHS TO \$100K ARR

SEO content compounding. Therapist partnership program. PR and podcast appearances. Each path is independent. The failure of any one does not prevent the business from working. The probability table in section 23 puts the combined probability of reaching \$100K at roughly 78% if any 2 channels fire and 88% if 3+ channels fire.

What success looks like at month 24

By May 2028, Mull Over has somewhere between 400 and 14,000 paying customers across NZ, AU, UK, Canada, and the US. Annual run-rate revenue is between \$30K and \$1.2M depending on which scenario plays out. The brand is recognisable in at least one of the five launch markets. There are 200 to 2,500 active therapist partners. The synthesis quality has improved through 50+ prompt iterations. Couples have completed somewhere between 2,000 and 80,000 shared decisions.

In all three scenarios the company is profitable. The founder is paid (in two of three scenarios). Optionality is wide: continue growing solo, raise capital and accelerate, or sell to an incumbent at a clean multiple.

The first action

Week one, day one: provision the Hetzner box, point mullover.ai DNS at it, push the static marketing site, invite the first beta couple to express interest by email. The clock starts there.

FINAL WORD

This is a buildable product, in a real market, with credible differentiation, on infrastructure that costs less per month than dinner out. The combination of low downside and meaningful upside is rare. The right answer is to build it, ship the MVP in 14 weeks, let the first 100 couples tell you whether the synthesis is as good as you think it is, and adjust from there.

End of business case, edition two